

INTERN PROJECT REPORT

DS_INT_SANDESH / Surprise Housing Case Study

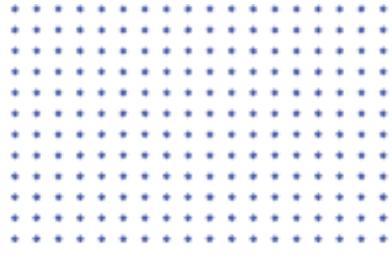




TABLE OF CONTENTS

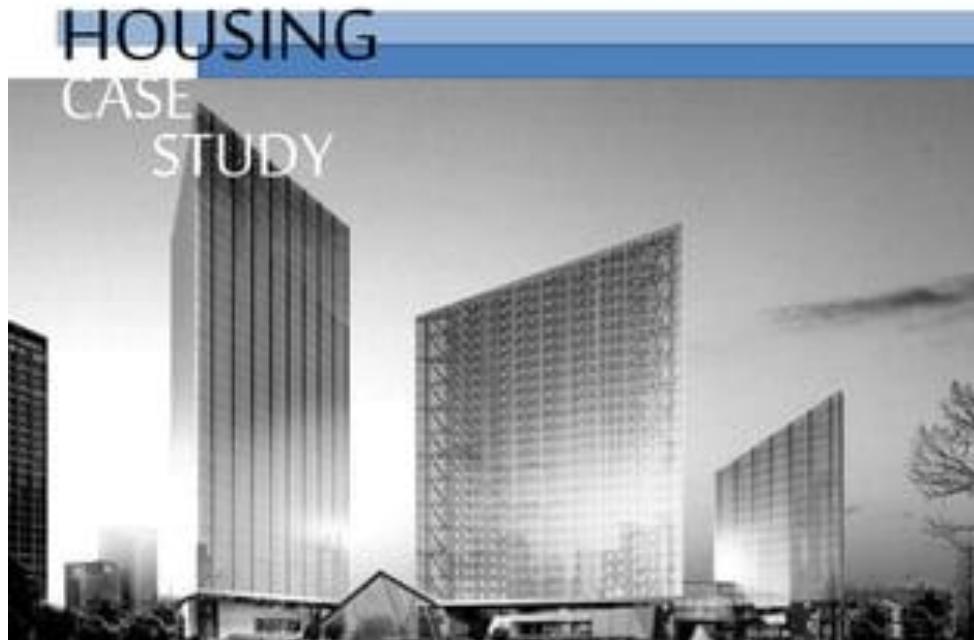
1. Abstract
2. Objective
3. Methodology
4. Code
5. Conclusion

ABSTRACT

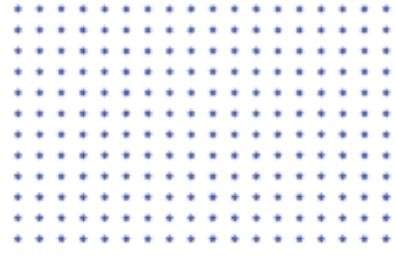


The "Surprise Housing" case study examines a unique residential development project aimed at addressing the critical issue of affordable housing in rapidly growing urban areas. This case study delves into the innovative strategies employed in the design, funding, and implementation of the Surprise Housing initiative. By analyzing the project's integration of sustainable practices, community involvement, and public-private partnerships, the study provides insights into its success factors and challenges. Key elements include the adaptation of modular construction techniques, use of green technologies, and the establishment of a robust stakeholder engagement process. The findings highlight how such a comprehensive approach can serve as a model for similar initiatives, offering valuable lessons for policymakers, developers, and community leaders in their efforts to create sustainable and inclusive housing solutions. The case study also identifies areas for improvement and future research to enhance the effectiveness of affordable housing projects.

Housing Case Study

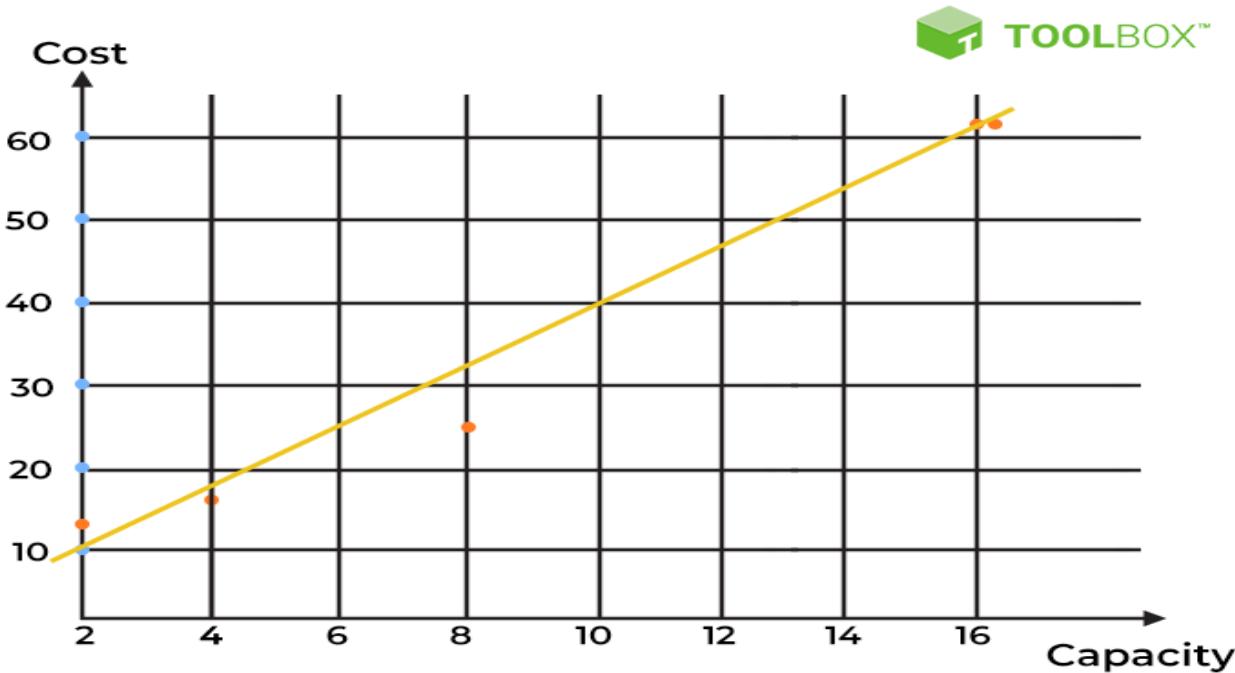


OBJECTIVE

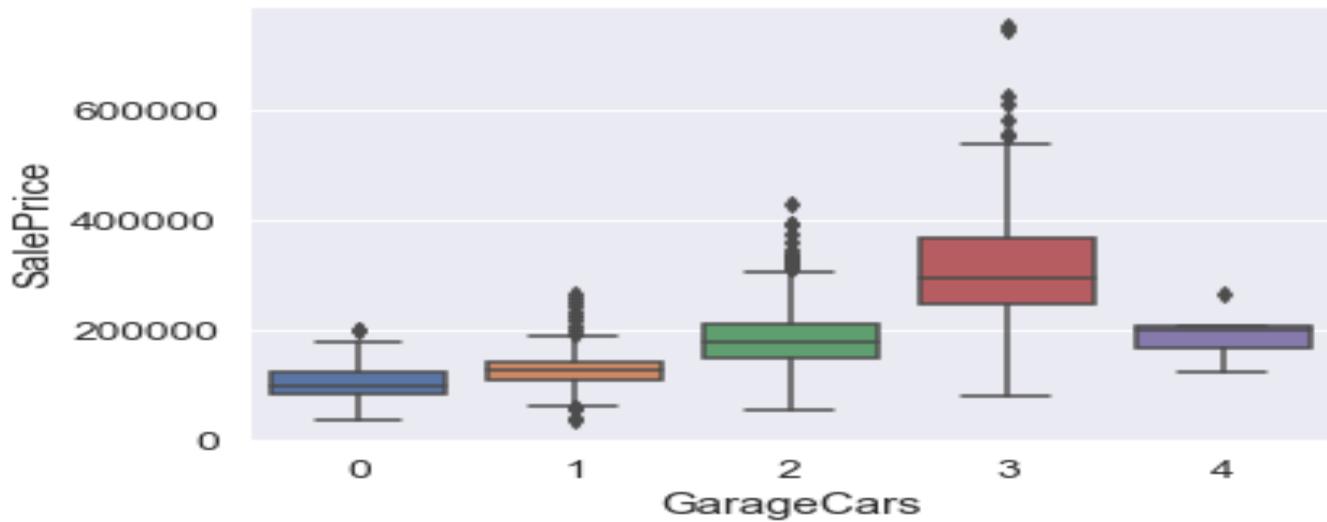


Objective

- **Determine the Influence of Key Variables:** Assess how different independent variables—such as construction costs, design features, and stakeholder engagement—affect dependent variables such as affordability, resident satisfaction, and sustainability.
- **Identify Significant Predictors:** Identify which factors are statistically significant predictors of project success metrics, including affordability and resident satisfaction, and quantify their effects.
- **Analyze Relationships between Variables:** Explore the relationships between independent variables and project outcomes to understand how they interact and contribute to overall project effectiveness.
- **Evaluate Model Fit and Robustness:** Assess the goodness-of-fit of the regression models to ensure they accurately represent the relationships between variables and provide reliable results.



METHODOLOGY



Data Preparation:

- **Data Cleaning:** Review and clean data to handle missing values, outliers, and inconsistencies.
- **Data Transformation:** Standardize variables and perform necessary transformations (e.g., log transformations for skewed data) to prepare for regression analysis.

Regression Analysis:

- **Model Selection:**
 - **Multiple Linear Regression:** Use multiple linear regression to analyze the relationship between independent variables (e.g., construction costs, design features) and dependent variables (e.g., affordability, resident satisfaction).
 - **Logistic Regression:** Apply logistic regression if analyzing categorical outcomes, such as the likelihood of meeting sustainability criteria.

CODE

Surprise Housing Case

localhost:8888/notebooks/Surprise%20Housing%20Case.ipynb

jupyter Surprise Housing Case Last Checkpoint: 2 days ago

File Edit View Run Kernel Settings Help Trusted

Code JupyterLab Python 3 (ipykernel)

	2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
...	
1455	1456	60	RL	62.0	7917	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
1456	1457	20	RL	85.0	13175	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	MnPrv	NaN	0	
1457	1458	70	RL	66.0	9042	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	GdPrv	Shed	2500	
1458	1459	20	RL	68.0	9717	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
1459	1460	20	RL	75.0	9937	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	NaN	0	

1460 rows × 81 columns

```
[5]: pd.set_option("display.max_columns",None)
pd.set_option("display.max_rows",None)
```

```
[6]: housing
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	LotConfig	LandSlope	Neighborhood	Condition1	Con...
0	1	60	RL	65.0	8450	Pave	NaN	Reg		Lvl	AllPub	Inside	Gtl	CollCr	Norm
1	2	20	RL	80.0	9600	Pave	NaN	Reg		Lvl	AllPub	FR2	Gtl	Veenker	Feedr
2	3	60	RL	68.0	11250	Pave	NaN	IR1		Lvl	AllPub	Inside	Gtl	CollCr	Norm

94°F Partly sunny

Search

ENG US 14:12 30-07-2024

Surprise Housing Case

localhost:8888/notebooks/Surprise%20Housing%20Case.ipynb

jupyter Surprise Housing Case Last Checkpoint: 2 days ago

File Edit View Run Kernel Settings Help Trusted

Code JupyterLab Python 3 (ipykernel)

	2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub	Inside	Gtl	CollCr	Norm
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub	Corner	Gtl	Crawfor	Norm	
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub	FR2	Gtl	NoRidge	Norm	
5	6	50	RL	85.0	14115	Pave	NaN	IR1	Lvl	AllPub	Inside	Gtl	Mitchel	Norm	
6	7	20	RL	75.0	10084	Pave	NaN	Reg	Lvl	AllPub	Inside	Gtl	Somerst	Norm	
7	8	60	RL	NaN	10382	Pave	NaN	IR1	Lvl	AllPub	Corner	Gtl	NWAmes	PosN	
8	9	50	RM	51.0	6120	Pave	NaN	Reg	Lvl	AllPub	Inside	Gtl	OldTown	Artery	

```
[7]: housing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 81 columns):
 #   Column          Non-Null Count  Dtype  
 ---  -- 
 0   Id              1460 non-null    int64  
 1   MSSubClass      1460 non-null    int64  
 2   MSZoning        1460 non-null    object  
 3   LotFrontage     1201 non-null    float64
 4   LotArea         1460 non-null    int64  
 5   Street          1460 non-null    object  
 6   Alley            91 non-null     object  
 7   LotShape         1460 non-null    object  
 8   LandContour     1460 non-null    object  
 9   Utilities        1460 non-null    object  
 10  LotConfig        1460 non-null    object  
 11  LandSlope        1460 non-null    object  
 12  Neighborhood     1460 non-null    object  
 13  Condition1      1460 non-null    object
```

94°F Partly sunny

Search

ENG US 14:12 30-07-2024

Surprise Housing Case

localhost:8888/notebooks/Surprise%20Housing%20Case.ipynb

jupyter Surprise Housing Case Last Checkpoint: 2 days ago

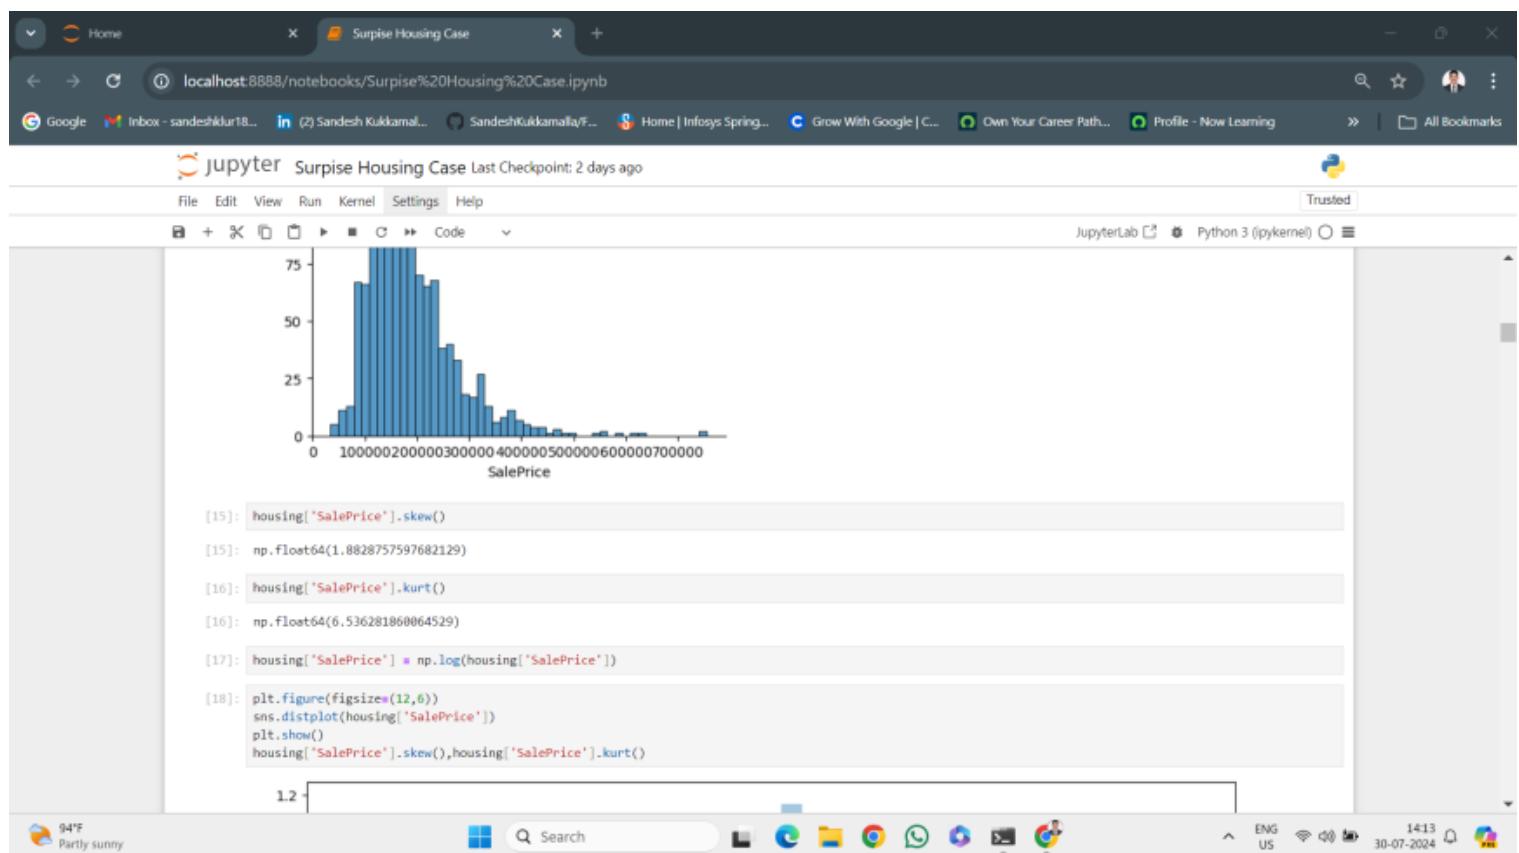
File Edit View Run Kernel Settings Help Trusted JupyterLab Python 3 (ipykernel)

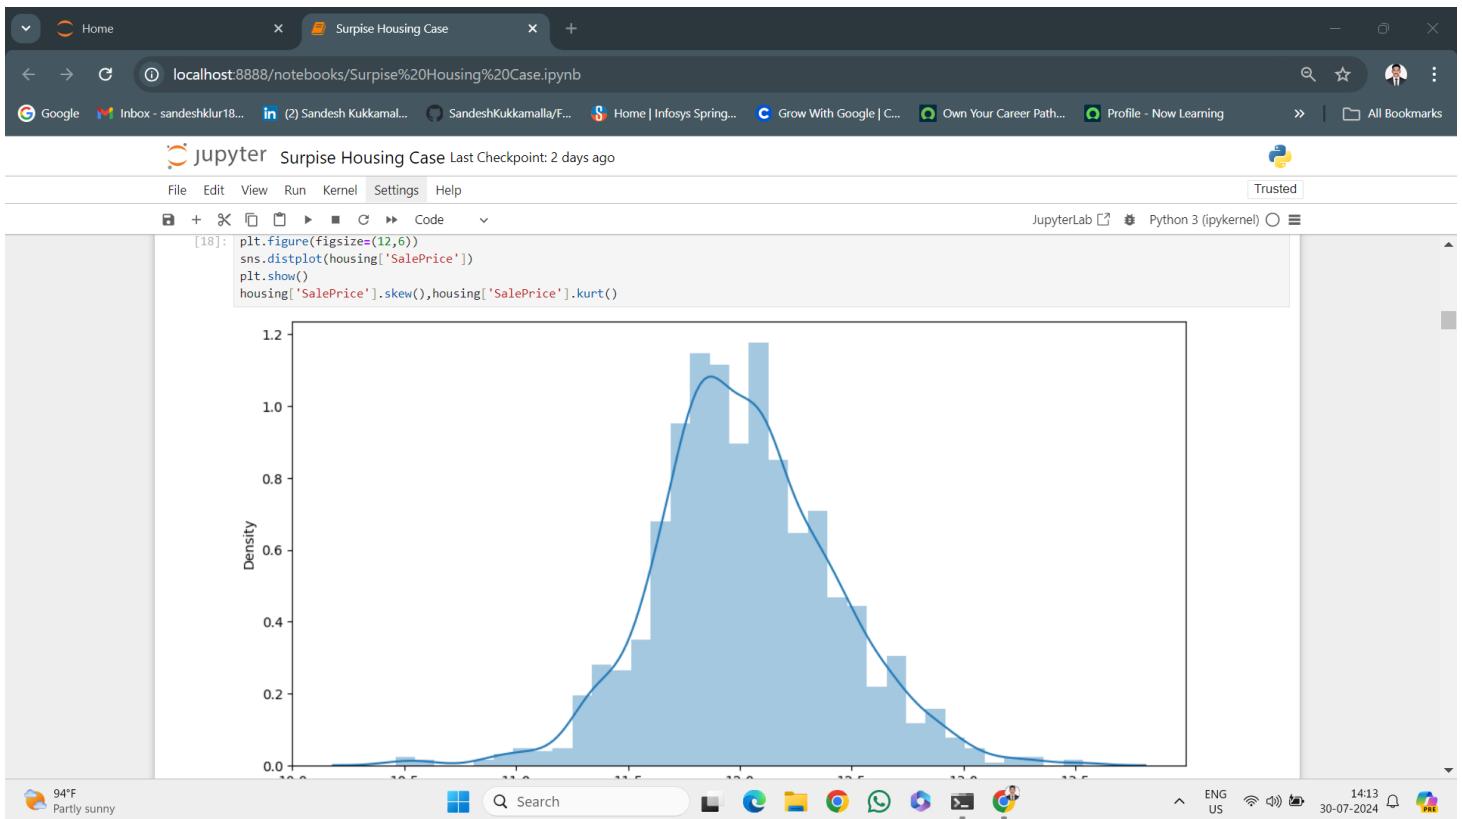
```
[8]: housing.isnull().mean()*100
```

```
[8]:
```

Id	0.00000	
MSSubClass	0.00000	
MSZoning	0.00000	
LotFrontage	17.739726	
LotArea	0.00000	
Street	0.00000	
Alley	93.767123	
LotShape	0.00000	
LandContour	0.00000	
Utilities	0.00000	
LotConfig	0.00000	
LandSlope	0.00000	
Neighborhood	0.00000	

94°F Partly sunny 1412 30-07-2024 ENG US





The screenshot shows a Jupyter Notebook interface with the title "Surprise Housing Case" and a subtitle "Last Checkpoint: 2 days ago". The notebook is running in a browser window with the URL "localhost:8888/notebooks/Surprise%20Housing%20Case.ipynb". The code cells [18] through [24] contain the following Python code:

```
[18]: (np.float64(0.12133506220520406), np.float64(0.8095319958036296))

[19]: housing.drop("Id", axis = 1, inplace = True)

[20]: housing[['MSubClass','OverallQual','OverallCond']] = \
    housing[['MSubClass','OverallQual','OverallCond']].astype('object')

[21]: #housing['LotFrontage'] = pd.to_numeric(housing ['LotFrontage'], errors = 'coerce')
#housing['MasVnrArea'] = pd.to_numeric(housing ['MasVnrArea'], errors = 'coerce')

[22]: #Filling Imputing collues with median/mean and create cols with Mode
null_cols = housing.columns[housing.isnull().any()]
null_cols

[22]: Index(['LotFrontage', 'MasVnrType', 'MasVnrArea', 'BsmtFinType1', 'Electrical',
       'GarageYrBlt'],
       dtype='object')

[23]: for feature in null_cols:
    if housing[feature].dtype == np.float64 or housing [feature].dtype == np.int64:
        housing[feature].fillna(housing[feature].median(),inplace=True)
    else:
        housing[feature].fillna(housing[feature].mode()[0],inplace=True)

[24]: housing.isnull().any().sum()

[24]: np.int64(0)
```

Surprise Housing Case

localhost:8888/notebooks/Surprise%20Housing%20Case.ipynb

jupyter Surprise Housing Case Last Checkpoint: 2 days ago

File Edit View Run Kernel Settings Help Trusted JupyterLab Python 3 (ipykernel)

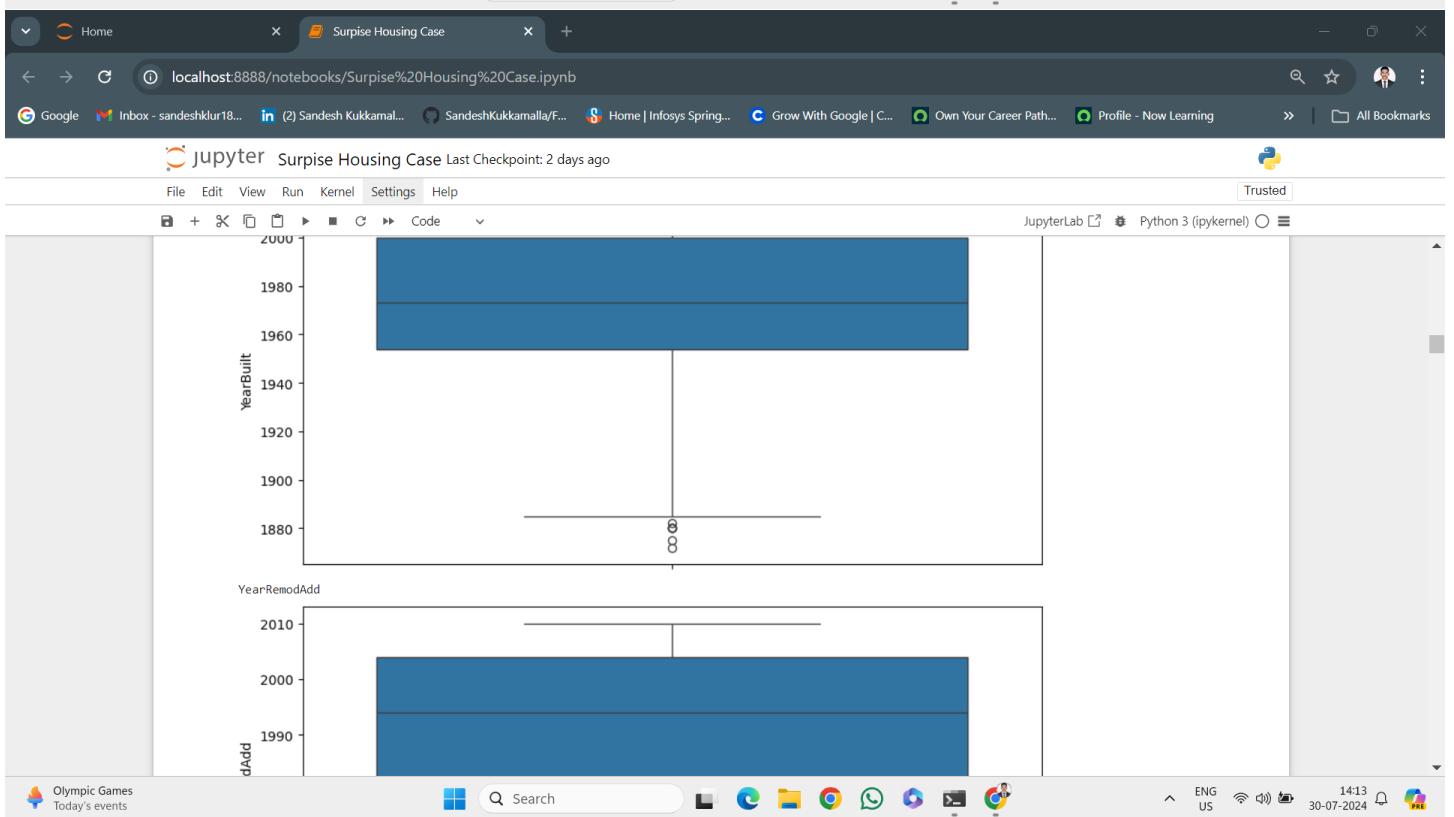
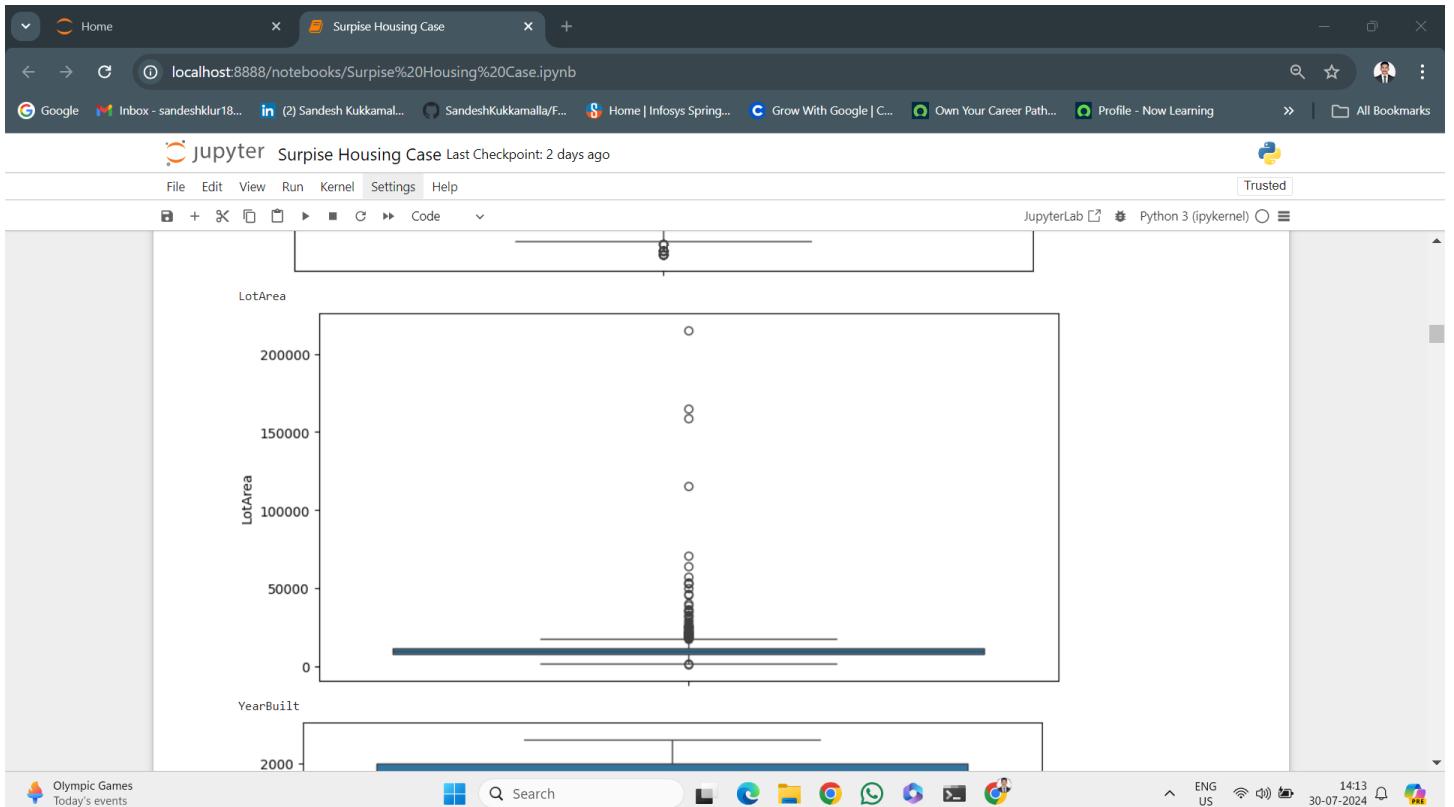
```
[29]: cat_cols = housing.select_dtypes(include = 'object').columns
num_cols = housing.select_dtypes(include = ['int64','float64']).columns
num_cols,cat_cols
```

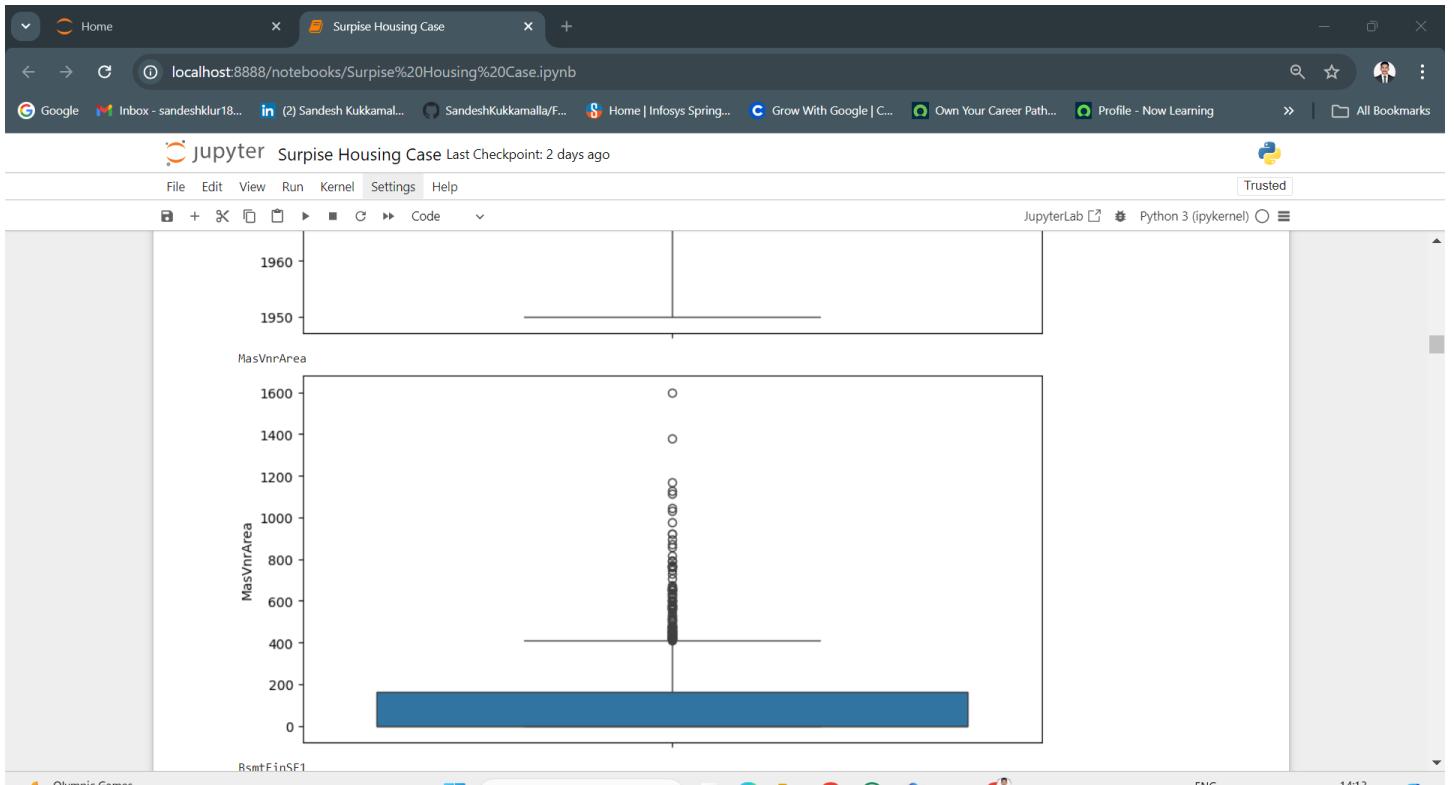
```
[29]: (Index(['LotFrontage', 'LotArea', 'YearBuilt', 'YearRemodAdd', 'MasVnrArea',
       'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfsF', 'TotalBsmtSF', '1stFlrSF',
       '2ndFlrSF', 'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmthalfBath',
       'FullBath', 'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'TotRmsAbvGr',
       'Fireplaces', 'GarageYrBlt', 'GarageCars', 'GarageArea', 'WoodDeckSF',
       'OpenPorchSF', 'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea',
       'MiscVal', 'MoSold', 'YrSold', 'SalePrice'],
      dtype='object'),
 Index(['MSubClass', 'MSZoning', 'Street', 'Alley', 'LotShape', 'LandContour',
       'Utilities', 'LotConfig', 'LandSlope', 'Neighborhood', 'Condition1',
       'Condition2', 'BldgType', 'HouseStyle', 'OverallQual', 'OverallCond',
       'RoofStyle', 'RoofMatl', 'Exteriorist', 'Exterior2nd', 'MasVnrType',
       'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual', 'BsmtCond',
       'BsmtExposure', 'BsmtFinType1', 'BsmtFinType2', 'Heating', 'HeatingQC',
       'CentralAir', 'Electrical', 'KitchenQual', 'Functional', 'FireplaceQu',
       'GarageType', 'GarageFinish', 'GarageQual', 'GarageCond', 'PavedDrive',
       'PoolQC', 'Fence', 'MiscFeature', 'SaleType', 'SaleCondition'],
      dtype='object'))
```

```
[38]: for col in num_cols:
    plt.figure(figsize=(10,5))
    print(col)
    sns.boxplot(housing[col])
    plt.show()
```

LotFrontage

LotArea

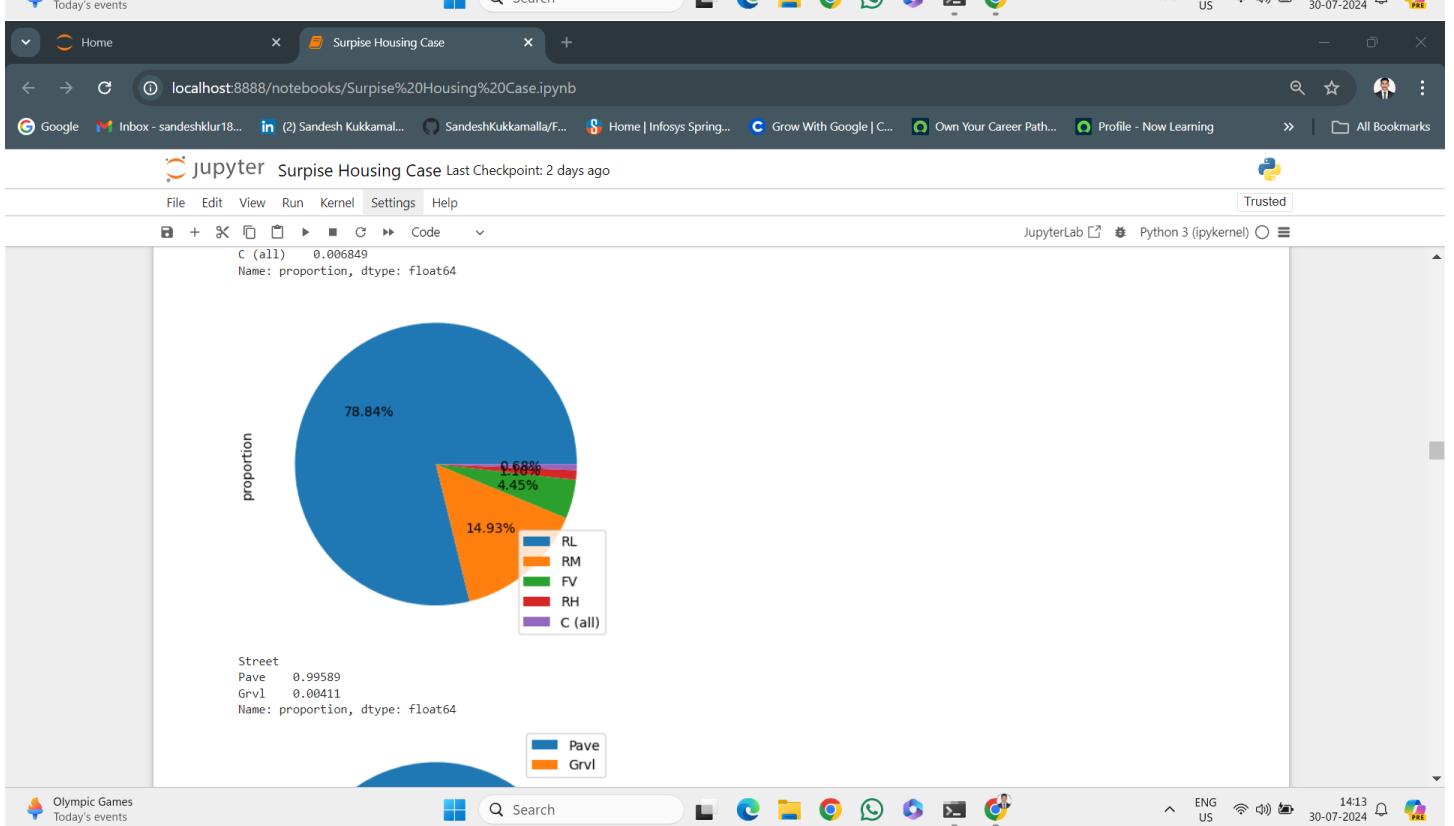
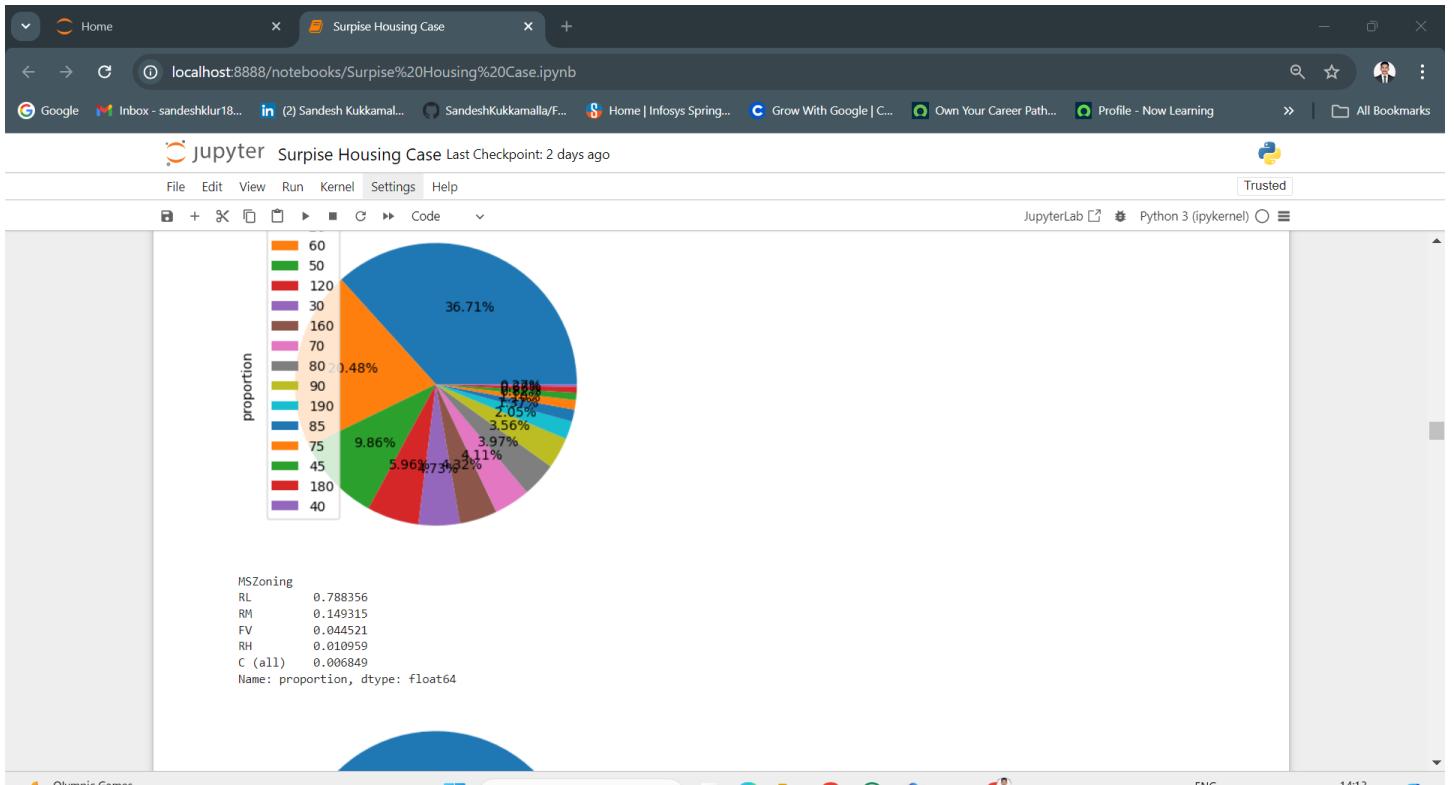


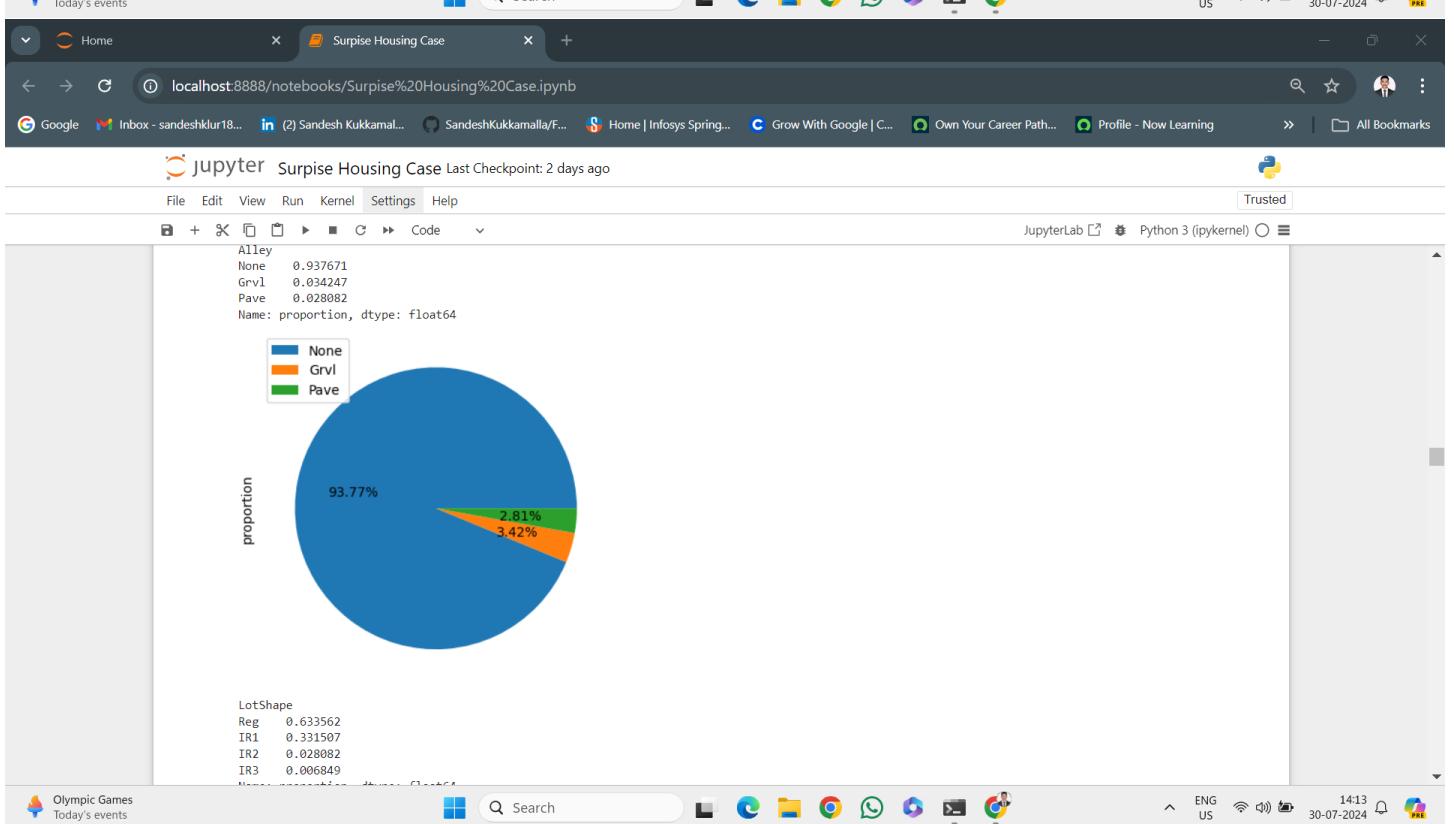
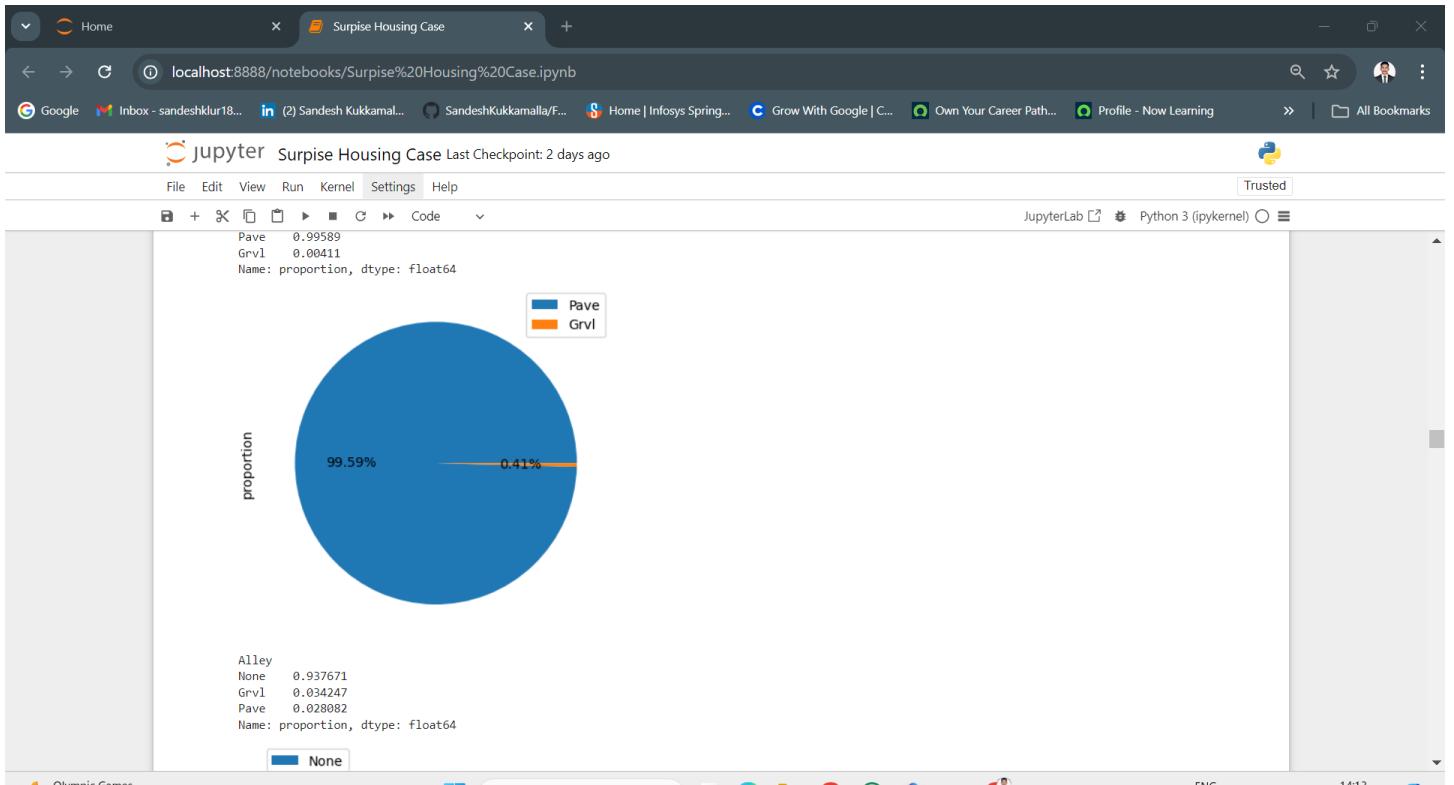


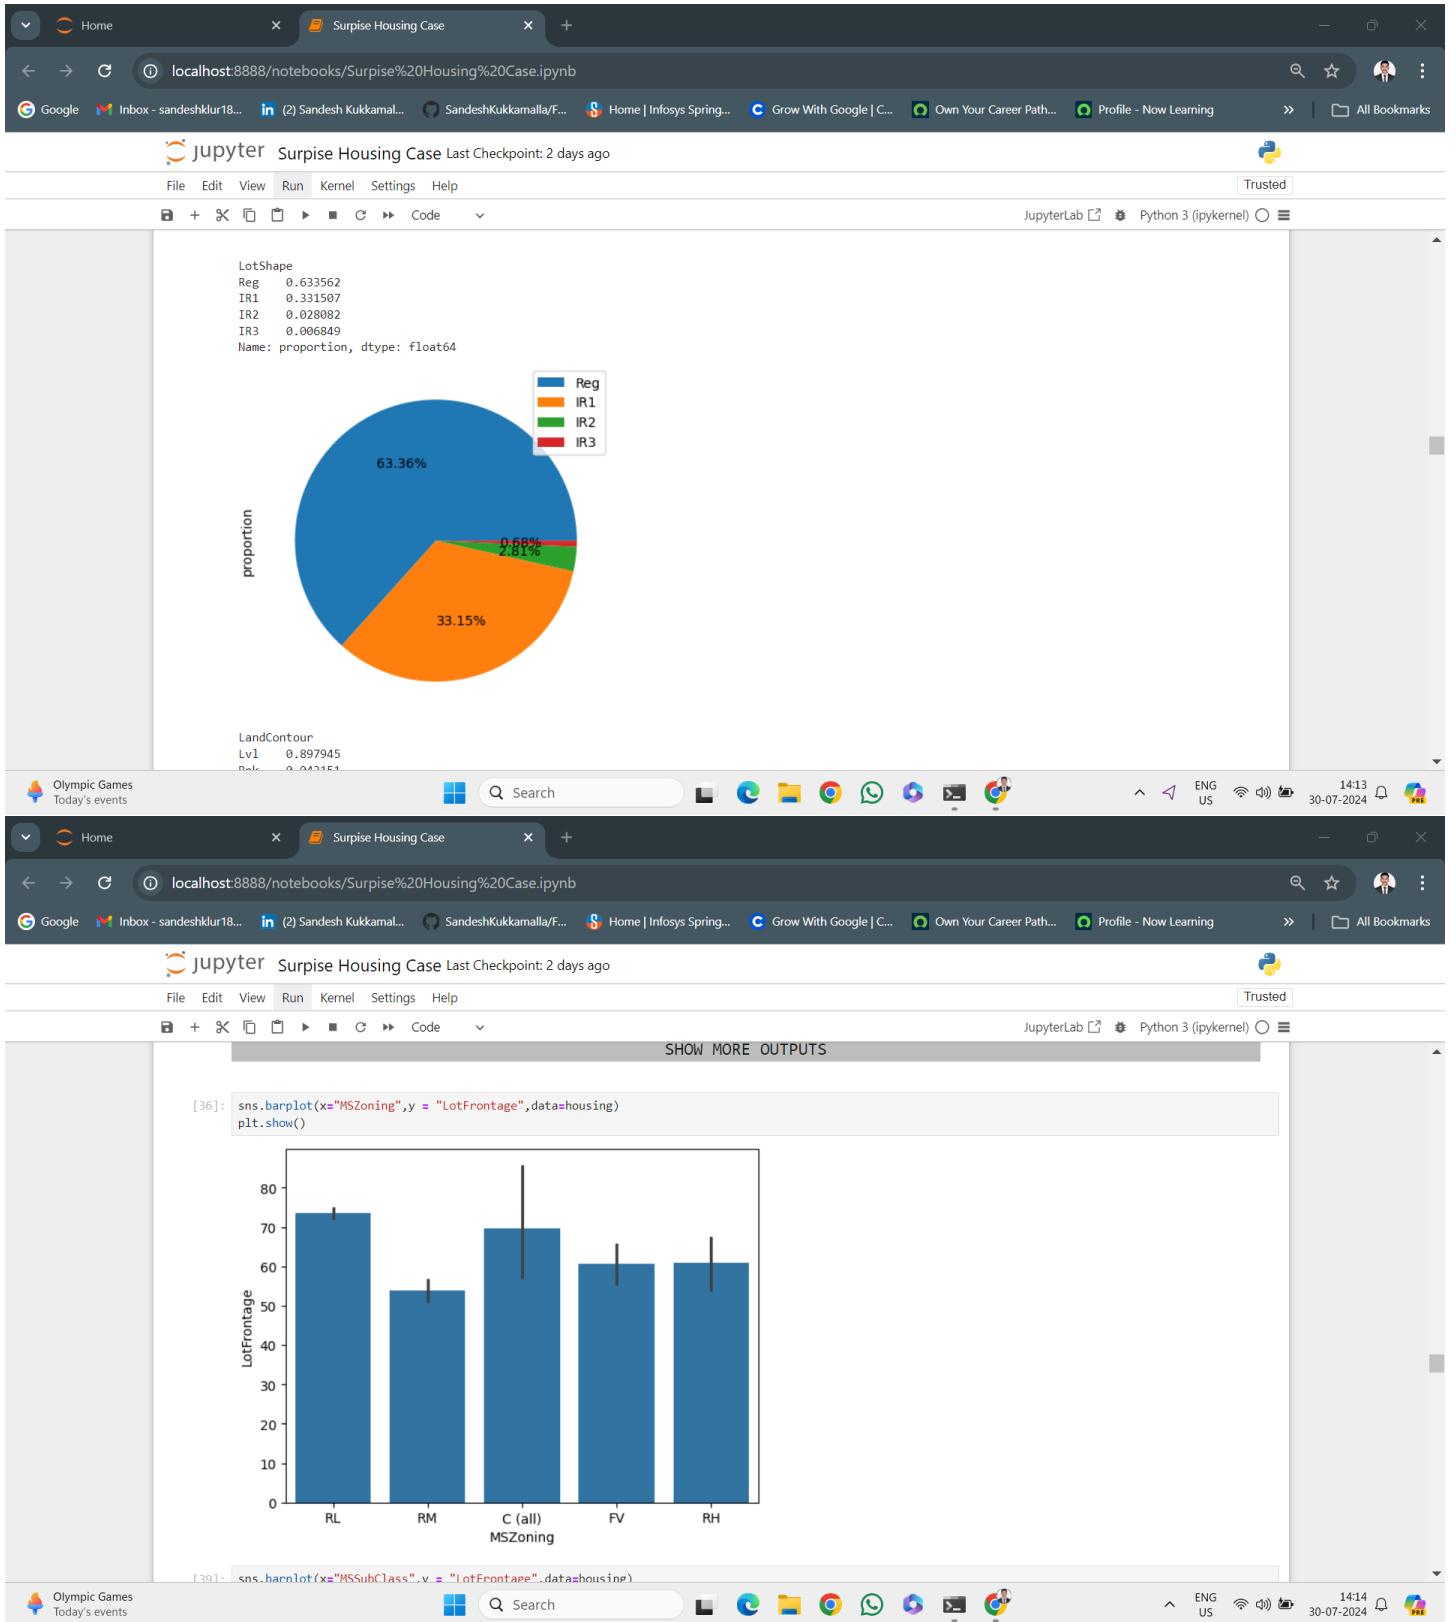
Surprise Housing Case Last Checkpoint: 2 days ago

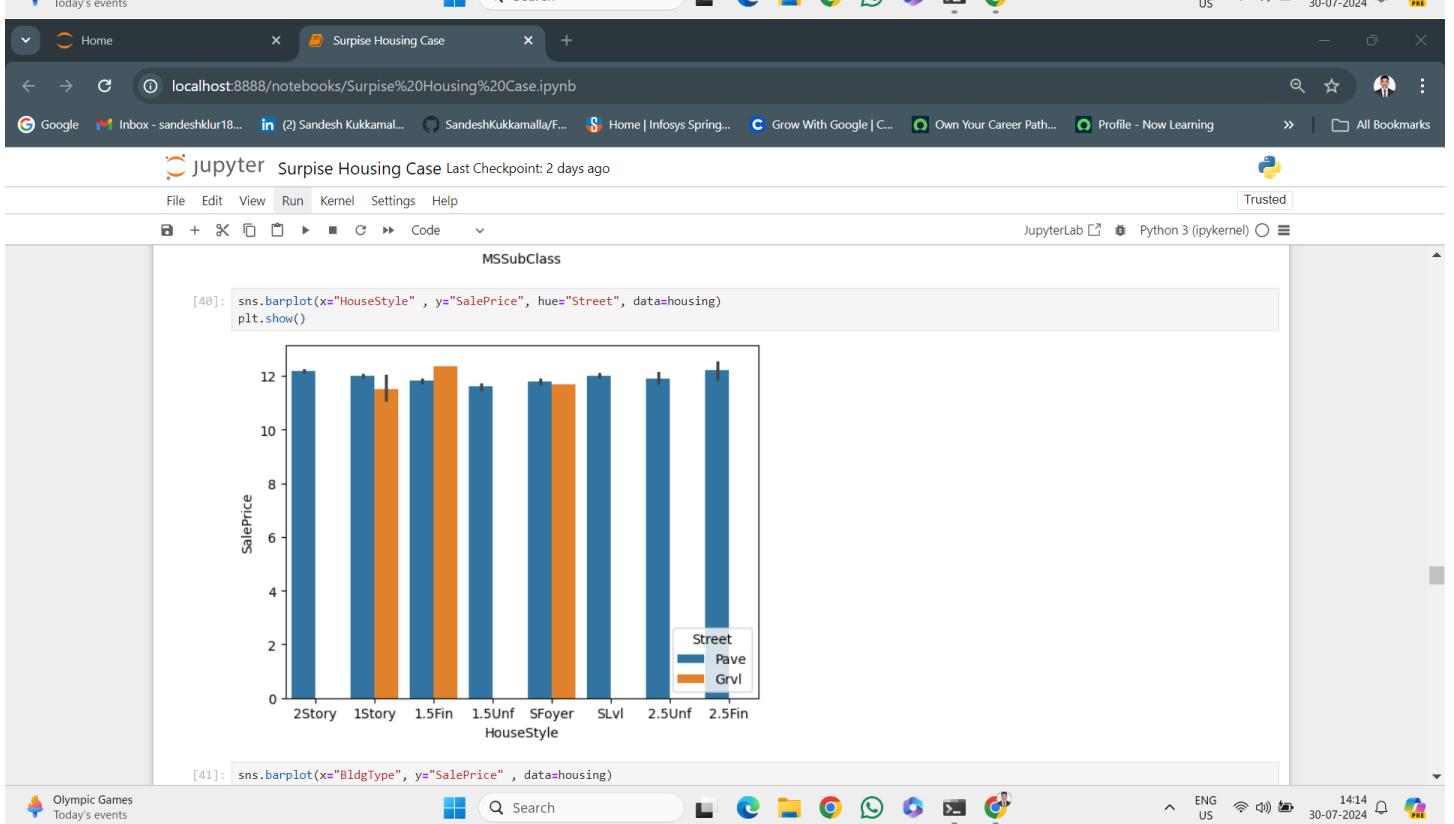
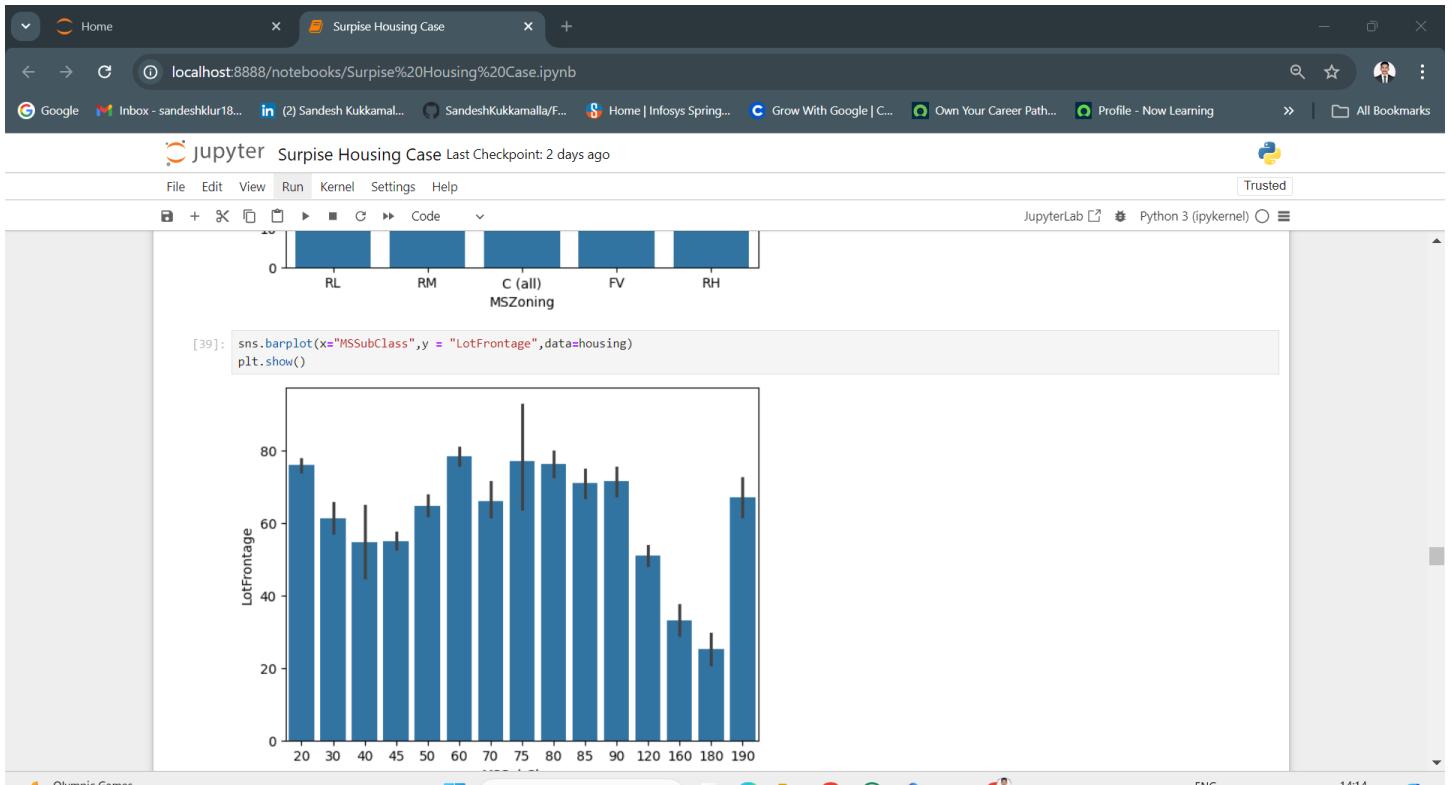
```
[32]: for col in cat_cols:
    print(housing[col].value_counts(normalize=True))
    plt.figure(figsize=(8,8))
    housing[col].value_counts(normalize=True).plot.pie(labeldistance=None, autopct="%1.2f%%")
    plt.legend()
    plt.show()
```

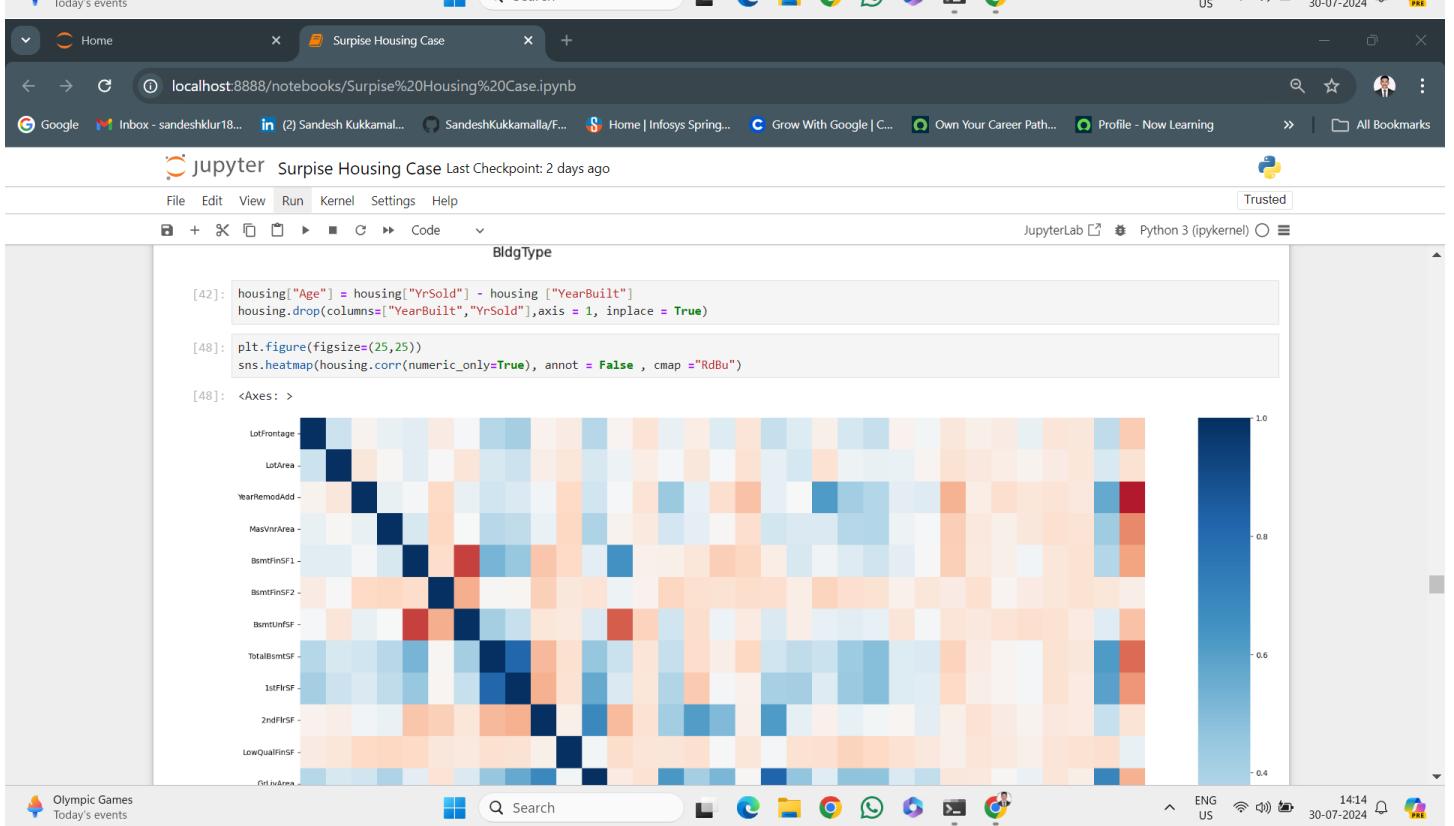
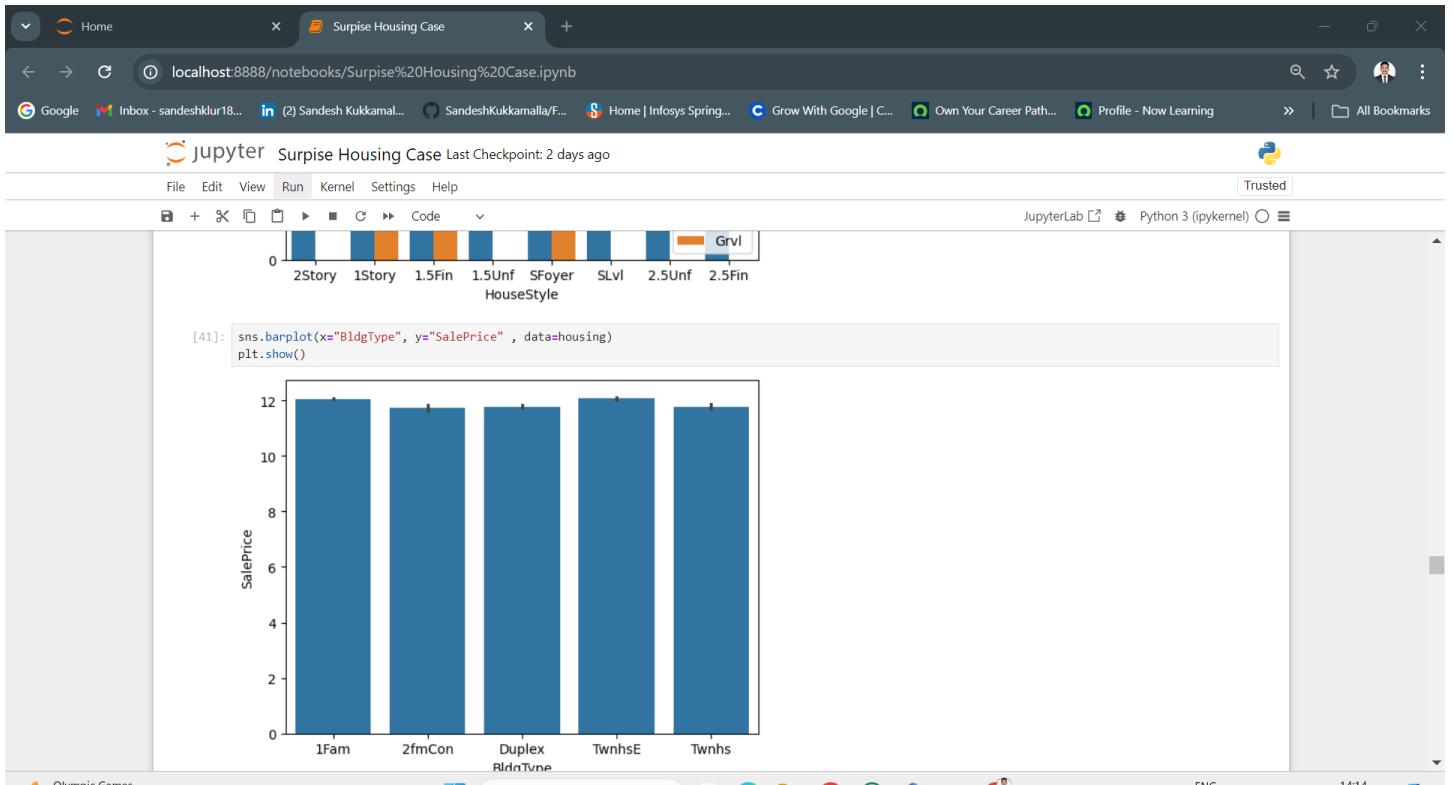
MSSubClass	proportion
20	0.367123
60	0.204795
50	0.098630
120	0.059589
30	0.047260
150	0.042151
70	0.041096
80	0.039726
90	0.035616
190	0.020548
85	0.013699
75	0.010959
45	0.008219
180	0.006849
40	0.002740

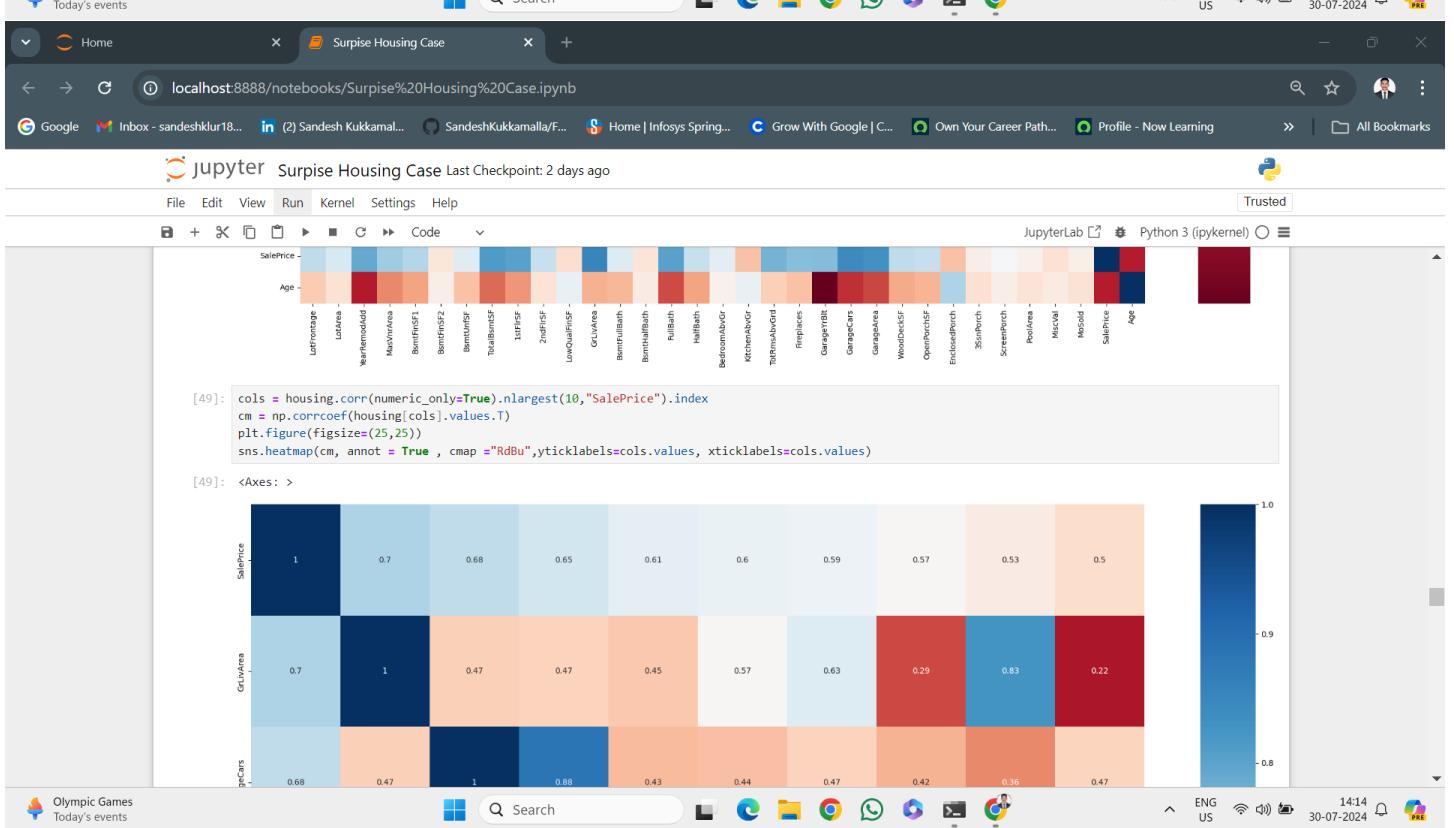
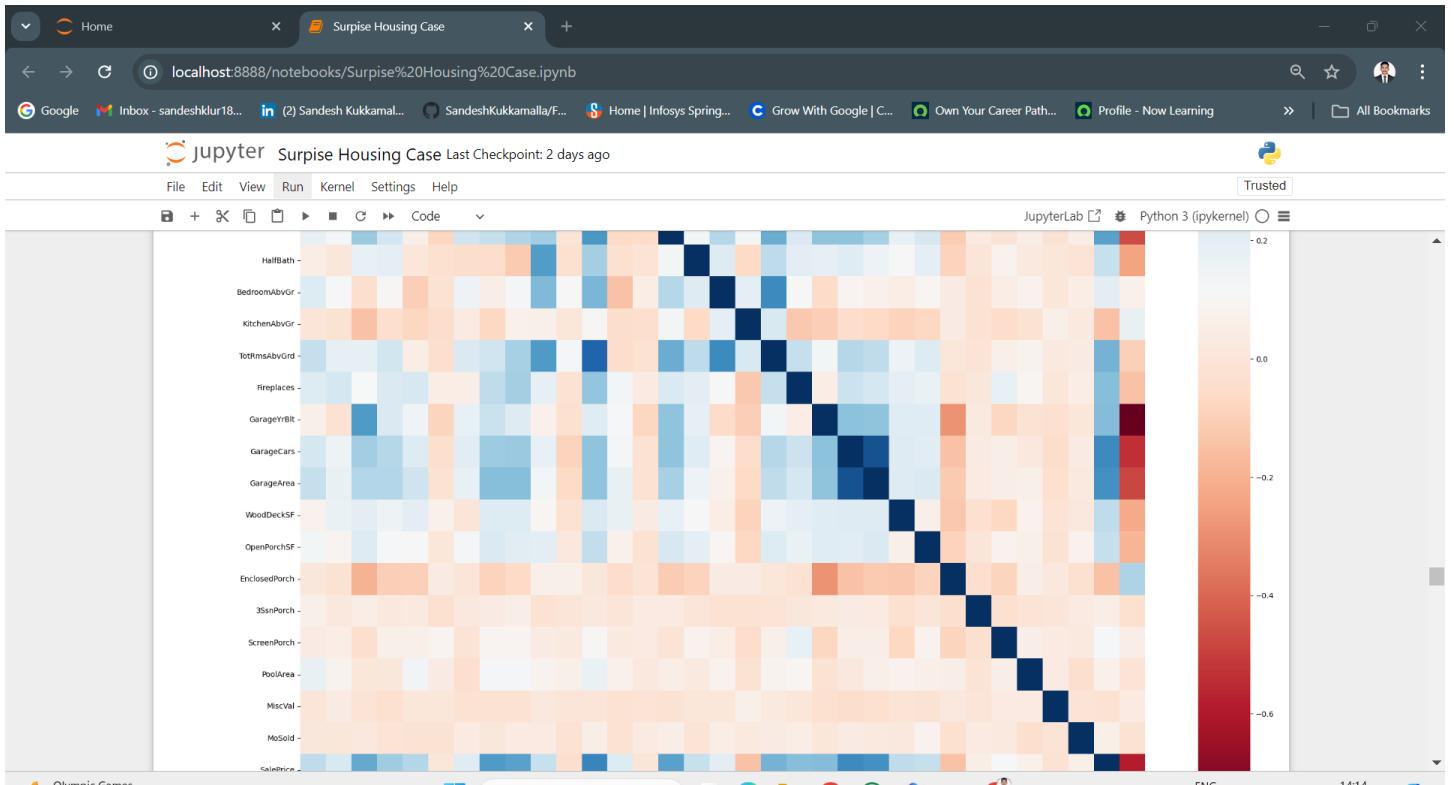


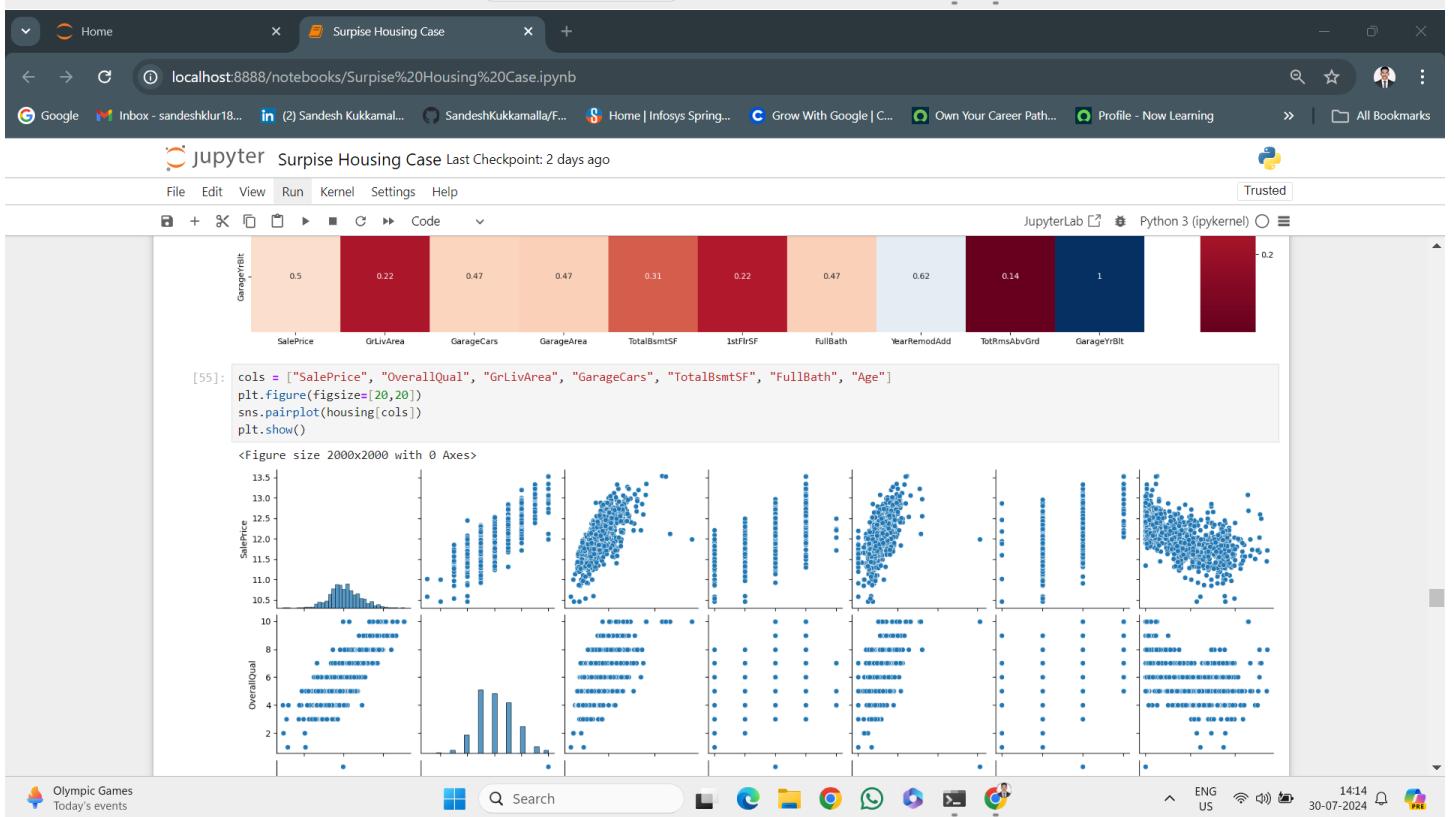
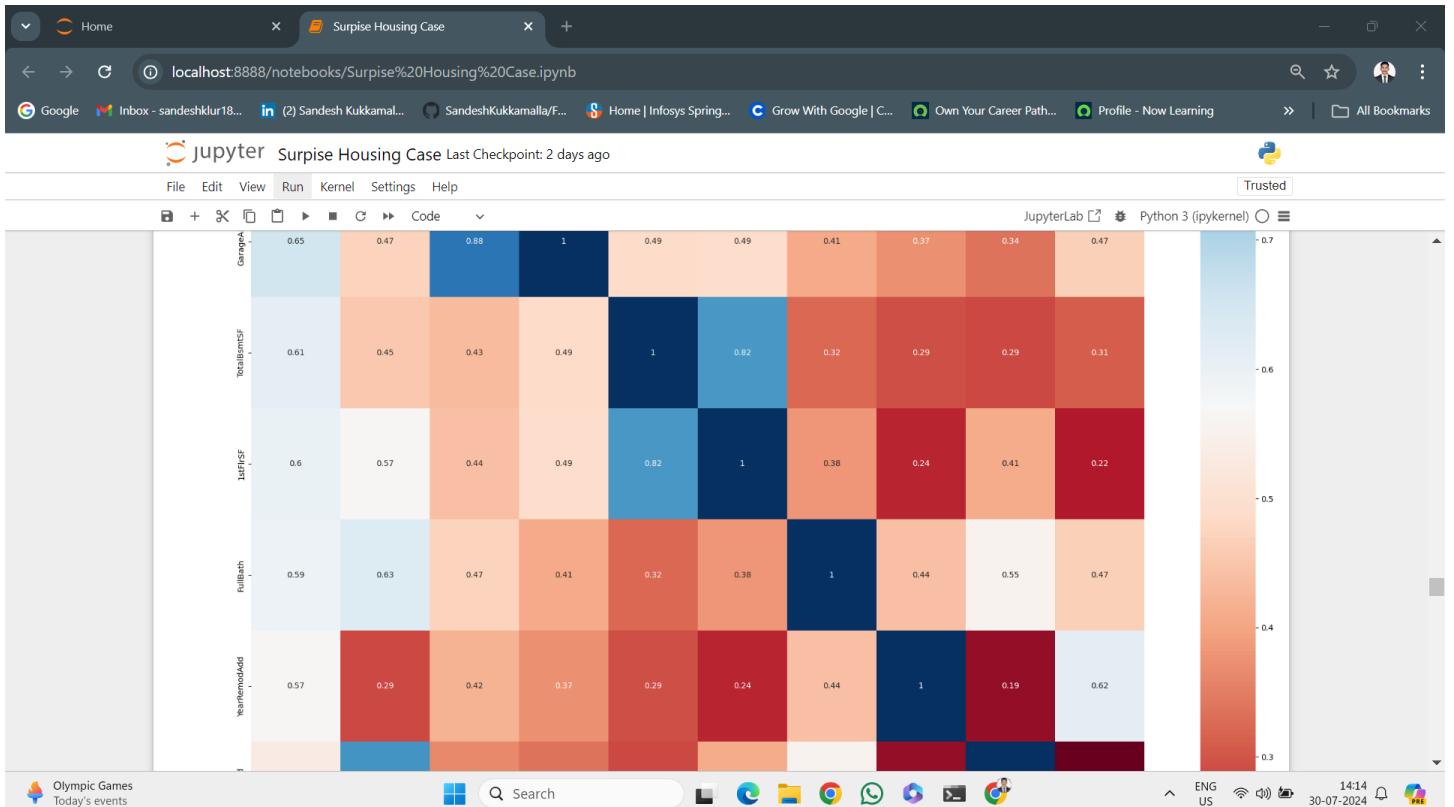


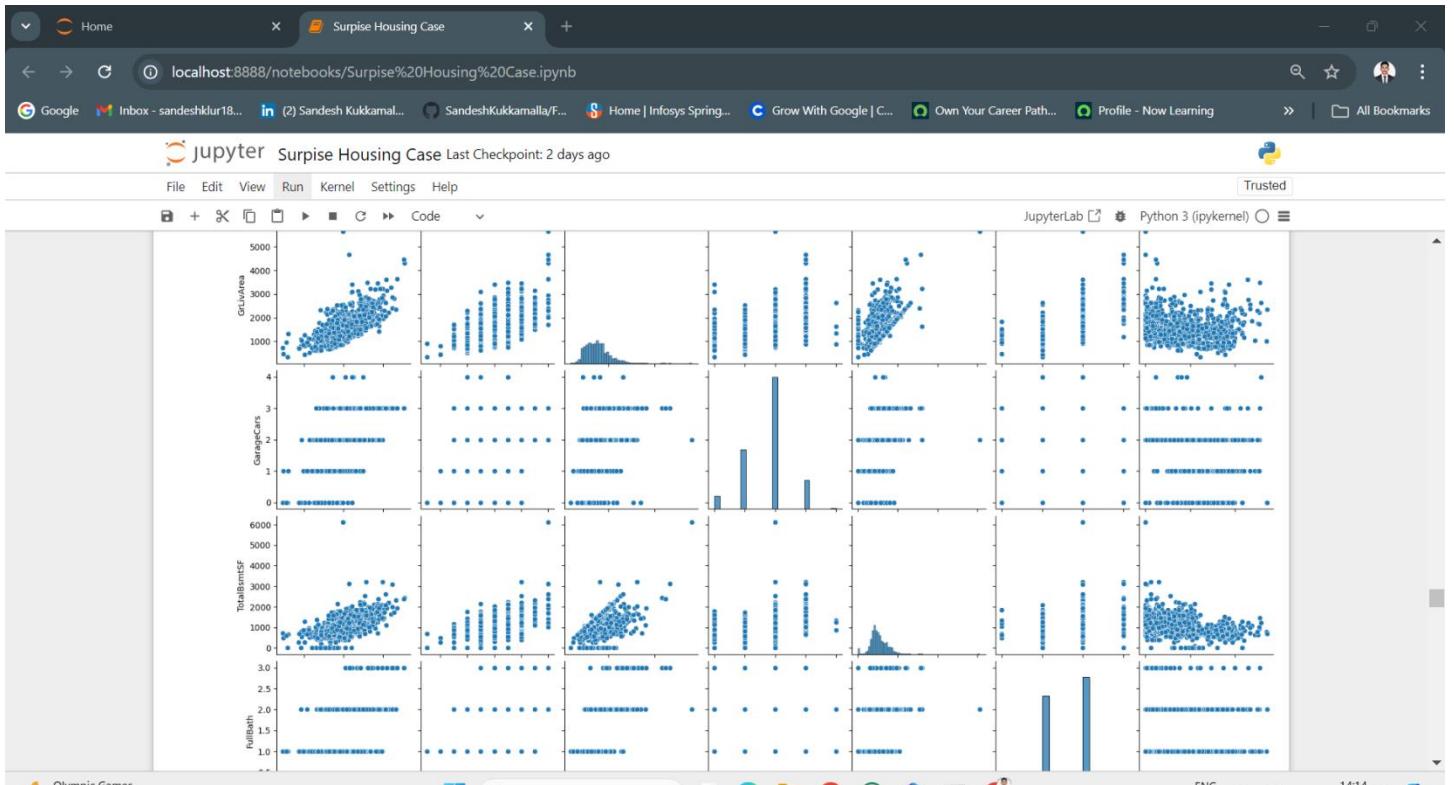












jupyter Surprise Housing Case Last Checkpoint: 2 days ago

```

File Edit View Run Kernel Settings Help Trusted JupyterLab Python 3 (ipykernel)

```

```

[58]: housing_num = housing.select_dtypes(include=['int64','float64'])
housing_cat = housing.select_dtypes(include='object')

[59]: housing_cat.shape
[60]: (1460, 46)

[61]: housing_cat_dm = pd.get_dummies(housing_cat,drop_first=True,dtype=int)

[62]: housing_cat_dm.shape
[63]: (1460, 252)

[64]: housing_cat_dm.head()

```

	MSSubClass_30	MSSubClass_40	MSSubClass_45	MSSubClass_50	MSSubClass_60	MSSubClass_70	MSSubClass_75	MSSubClass_80	MSSubClass_85	MSSubClass_90
0	0	0	0	0	0	1	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0
2	0	0	0	0	1	0	0	0	0	0
3	0	0	0	0	0	0	1	0	0	0

```
[64]: house = pd.concat([housing_cat_dm, housing_num], axis=1)
[65]: house.shape
[65]: (1460, 285)
[67]: X = house.drop(['SalePrice'], axis=1).copy()
y=house['SalePrice'].copy()
[68]: import sklearn
[69]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
[70]: X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.25,random_state=42)
[72]: num_cols = list(X_train.select_dtypes(include=['int64','float64']).columns)
[74]: scaler = StandardScaler()
[77]: X_train[num_cols]=scaler.fit_transform(X_train[num_cols])
X_test[num_cols] = scaler.fit_transform(X_test[num_cols])
[121]: def eval_metrics(y_train, y_train_pred, y_test ,y_pred):
    print("r2 score (train) = ",%2f % r2_score(y_train,y_train_pred))
```

```
[77]: X_train[num_cols]=scaler.fit_transform(X_train[num_cols])
X_test[num_cols] = scaler.fit_transform(X_test[num_cols])
[121]: def eval_metrics(y_train, y_train_pred, y_test ,y_pred):
    print("r2 score (train) = ",%2f % r2_score(y_train,y_train_pred))
    print("r2 score (test) = ",%2f % r2_score(y_test,y_pred))

    mse_train = mean_squared_error(y_train,y_train_pred)
    mse_test = mean_squared_error(y_test, y_pred)
    rmse_train = mse_train **0.5
    rmse_test = mse_test ** 0.5

    print("RMSE(Train) = ",%0.2f % rmse_train)
    print("RMSE(Test) = ",%0.2f% rmse_test)

[122]: import sklearn
from sklearn.feature_selection import RFE
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_squared_error
from sklearn.linear_model import Ridge,Lasso
from sklearn.model_selection import GridSearchCV
[123]: #Applying Ridge regressing with varing the hyperparameter 'Lamda'

params = { 'alpha':
           [0.0001,0.001,0.01,0.05,0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,1.0,2.0,3.0,4.0,5.0,6.0,7.0,8.0,9.0,10,20,50,100,500,1000]}
ridge = Ridge()
ridgeCV = GridSearchCV(estimator=ridge,param_grid=params,scoring='neg_mean_absolute_error',cv=5,
                       return_train_score=True, verbose=1, n_jobs=-1)
```

Surprise Housing Case

```
File Edit View Run Kernel Settings Help Trusted
[123]: params = {'alpha': [0.0001, 0.001, 0.01, 0.05, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 2.0, 3.0, 4.0, 5.0, 6.0, 7.0, 8.0, 9.0, 10, 20, 50, 100, 500, 1000]}
ridge = Ridge()
ridgeCV = GridSearchCV(estimator=ridge, param_grid=params, scoring='neg_mean_absolute_error', cv=5, return_train_score=True, verbose=1, n_jobs=-1)
ridgeCV.fit(X_train, y_train)

Fitting 5 folds for each of 28 candidates, totalling 140 fits
[123]: + GridSearchCV ⓘ ⓘ
  + best_estimator_: Ridge
    + Ridge ⓘ ⓘ
      Ridge(alpha=9)

[124]: ridgeCV.best_params_
[124]: {'alpha': 100}

[125]: ridge = Ridge(alpha=9)
ridge.fit(X_train, y_train)

[125]: + Ridge ⓘ ⓘ
  Ridge(alpha=9)

[145]: y_train_pred = ridge.predict(X_train)
y_pred = ridge.predict(X_test)
```

Surprise Housing Case

```
File Edit View Run Kernel Settings Help Trusted
[145]: Ridge(alpha=9)

[145]: y_train_pred = ridge.predict(X_train)
y_pred = ridge.predict(X_test)

[146]: eval_metrics(y_train, y_train_pred, y_test, y_pred)

r2 score (train) = 0.945534
r2 score (test) = 0.88
RMSE(Train) = 0.09
RMSE(test) = 0.2f

[147]: ridgeCV_res = pd.DataFrame(ridgeCV.cv_results_)
ridgeCV_res.head()

[147]: mean_fit_time std_fit_time mean_score_time std_score_time param_alpha params split0_test_score split1_test_score split2_test_score split3_test_score split4_test_s
      0 0.064471 0.029654 0.020168 0.010803 0.0001 {'alpha': 0.0001} -0.096378 -0.120607 -0.105644 -0.093268 -0.09
      1 0.058076 0.024525 0.017421 0.005274 0.0010 {'alpha': 0.01} -0.096377 -0.120556 -0.105629 -0.093263 -0.09
      2 0.048130 0.011092 0.023170 0.008887 0.0100 {'alpha': 0.01} -0.096363 -0.120099 -0.105489 -0.093221 -0.09
      3 0.054286 0.010736 0.018098 0.003221 0.0500 {'alpha': 0.05} -0.096312 -0.118790 -0.104979 -0.093063 -0.09
      4 0.063639 0.018677 0.023806 0.011471 0.1000 {'alpha': 0.1} -0.096262 -0.117923 -0.104484 -0.092892 -0.09
```

Surprise Housing Case

localhost:8888/notebooks/Surprise%20Housing%20Case.ipynb

jupyter Surprise Housing Case Last Checkpoint: 2 days ago

File Edit View Run Kernel Settings Help Trusted JupyterLab Python 3 (ipykernel)

	3	0.054286	0.010736	0.018098	0.003221	0.0500	{'alpha': 0.05}	-0.096312	-0.118790	-0.104979	-0.093063	-0.09
	4	0.063639	0.018677	0.023806	0.011471	0.1000	{'alpha': 0.1}	-0.096262	-0.117923	-0.104484	-0.092892	-0.09

```
[148]: plt.plot(ridgeCV_res['param_alpha'],ridgeCV_res['mean_train_score'],label ='train')
plt.plot(ridgeCV_res['param_alpha'],ridgeCV_res['mean_train_score'],label = 'test')
plt.xlabel('alpha')
plt.ylabel('R2_score')
plt.xscale('log')
plt.legend()
plt.show()
```

Olympic Games Today's events Search ENG US 14:15 30-07-2024

Surprise Housing Case

localhost:8888/notebooks/Surprise%20Housing%20Case.ipynb

jupyter Surprise Housing Case Last Checkpoint: 2 days ago

File Edit View Run Kernel Settings Help Trusted JupyterLab Python 3 (ipykernel)

```
[149]: params = { 'alpha':
[0.0001,0.001,0.01,0.05,0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,1.0,2.0,3.0,4.0,5.0,6.0,7.0,8.0,9.0,10,20,50,100,500,1000]}
lasso = Lasso()
lassoCV = GridSearchCV(estimator=ridge,param_grid=params,scoring='neg_mean_absolute_error',cv=5,
return_train_score=True, verbose=1, n_jobs=-1)
lassoCV.fit(X_train,y_train)
```

Fitting 5 folds for each of 28 candidates, totalling 140 fits

```
[149]: GridSearchCV
+ best_estimator_: Ridge
  + Ridge
```

```
[150]: lassoCV.best_params_
[150]: {'alpha': 100}
```

```
[151]: lasso=Lasso(alpha=0.001)
```

Olympic Games Today's events Search ENG US 14:15 30-07-2024

Surprise Housing Case

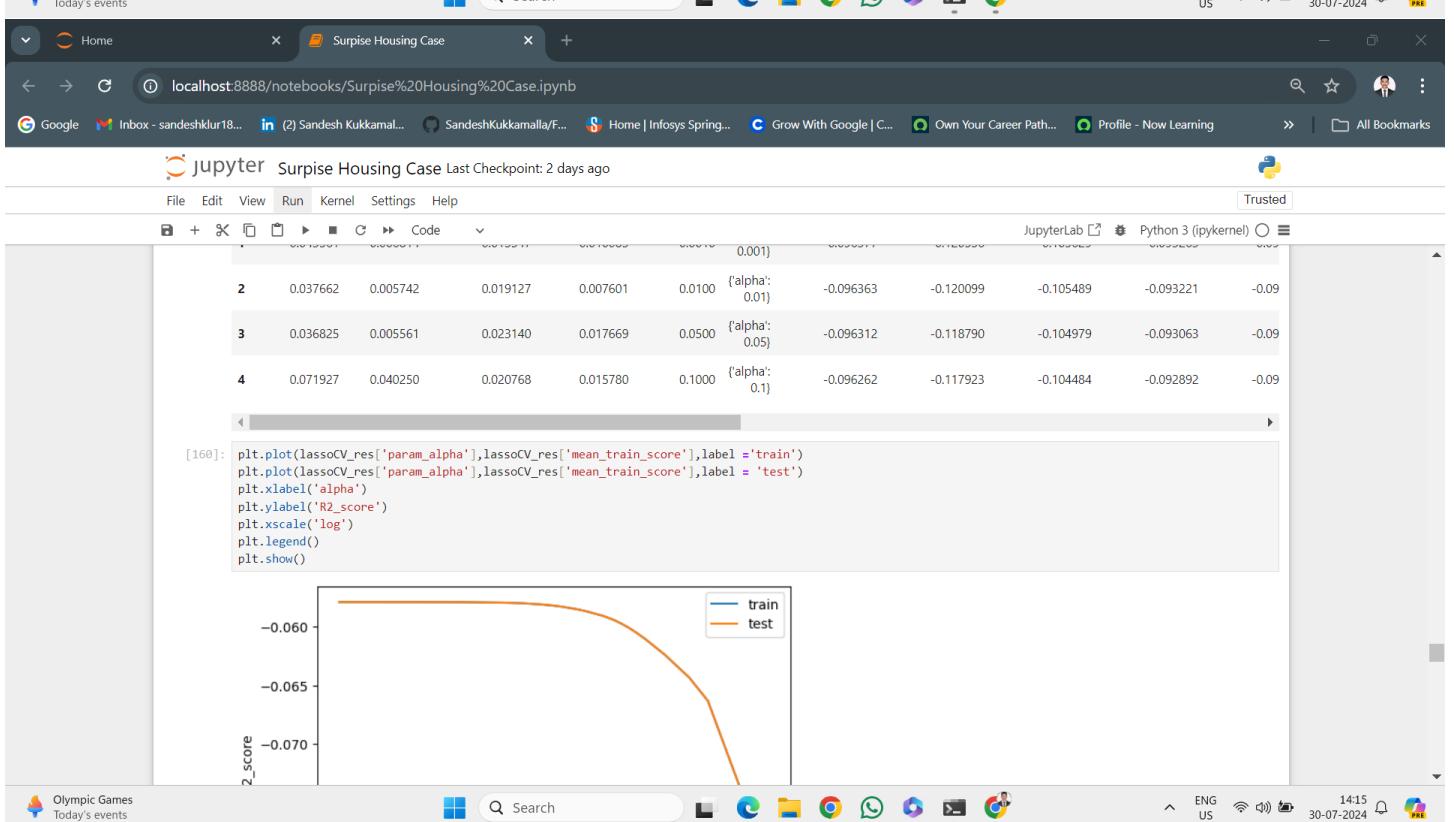
```
[150]: lassoCV.best_params_
[150]: {'alpha': 100}
[151]: lasso=Lasso(alpha=0.001)
lasso.fit(X_train,y_train)
[151]: Lasso(alpha=0.001)

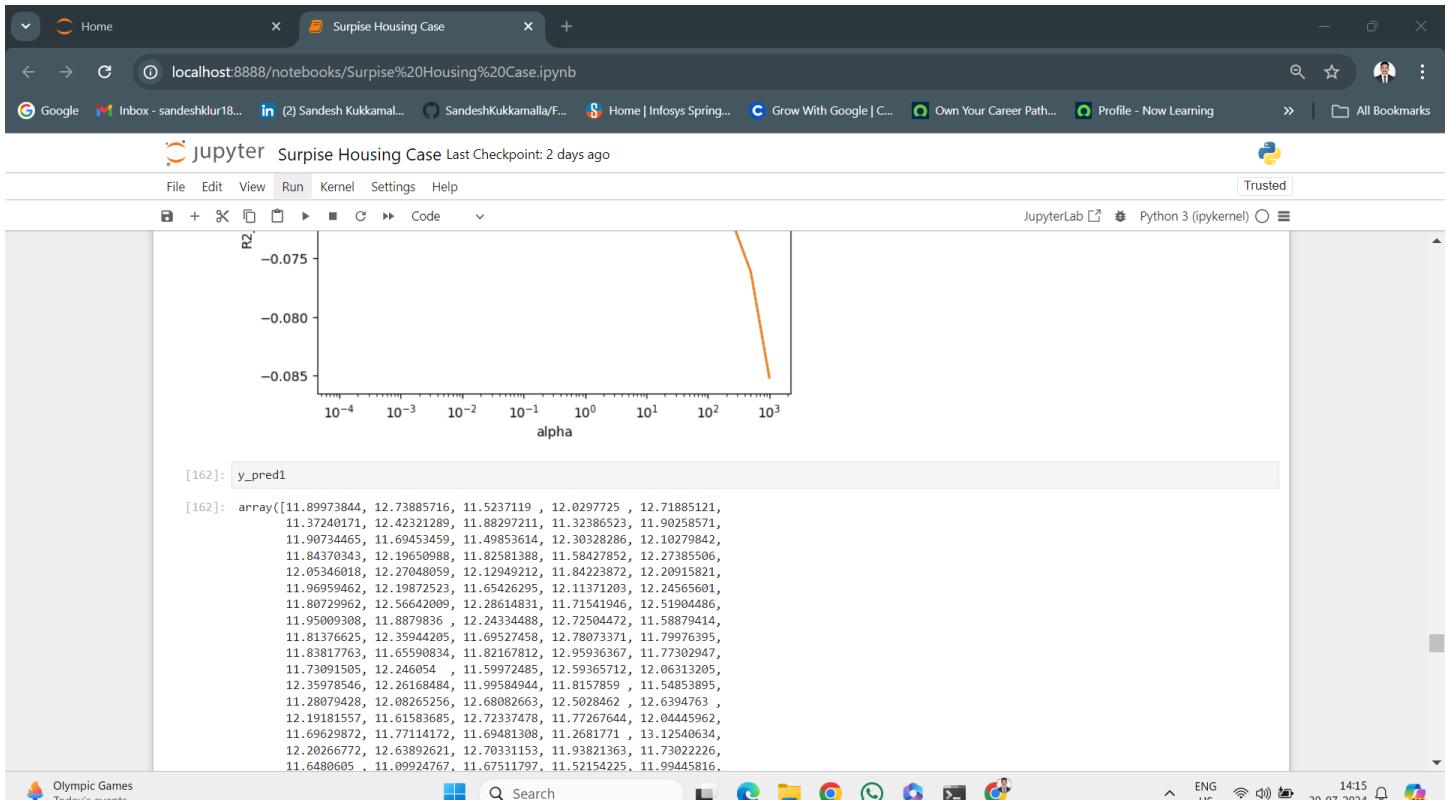
[157]: y_train_pred1 = lasso.predict(X_train)
y_pred1 = lasso.predict(X_test)

[158]: eval_metrics(y_train,y_train_pred1, y_test, y_pred1)
r2 score (train) =  0.939695
r2 score (test) = 0.88
RMSE(Train) =  0.10
RMSE(Test) =  0.26

[159]: lassoCV_res = pd.DataFrame(lassoCV.cv_results_)
lassoCV_res.head()

[159]:   mean_fit_time  std_fit_time  mean_score_time  std_score_time  param_alpha  params  split0_test_score  split1_test_score  split2_test_score  split3_test_score  split4_test_s
      0       0.052941     0.010117      0.012133      0.004352      0.0001  {'alpha': 0.0001}      -0.096378     -0.120607     -0.105644     -0.093268      -0.09
      1       0.043561     0.006814      0.015547      0.010085      0.0010  {'alpha': 0.001}      -0.096377     -0.120556     -0.105629     -0.093263      -0.09
```





The figure shows a Jupyter Notebook interface with the same title and last checkpoint as the first one. The menu bar and Trusted button are present. A code cell [165] contains the following Python code to calculate Ridge and Lasso coefficients:

```
betas = pd.DataFrame(index=X.columns)
betas.rows = X.columns
betas['Ridge'] = ridge.coef_
betas['Lasso'] = lasso.coef_
betas
```

Below the code cell is a table titled "Ridge" and "Lasso" showing the coefficients for various columns. The columns are labeled with MSSubClass values: 30, 40, 45, 50, 60, 70, 75, 80, 85, 90, and 120. The Ridge coefficients are mostly negative, while the Lasso coefficients are mostly zero or very small. The notebook interface includes a toolbar at the bottom with icons for file operations, search, and other common functions.

	MSSubClass_160	-0.024910	-1.989819e-02
MSSubClass_180	-0.002585	-1.081793e-03	
MSSubClass_190	-0.001511	-0.000000e+00	
MSZoning_FV	0.037276	1.076613e-02	
MSZoning_RH	0.016375	2.412842e-03	
MSZoning_RL	0.059178	1.079683e-02	
MSZoning_RM	0.046035	-0.000000e+00	
Street_Pave	0.002325	8.893126e-04	
Alley_None	-0.005913	-2.348874e-03	
Alley_Pave	0.007911	8.837248e-03	
LotShape_IR2	0.006331	5.848765e-03	
LotShape_IR3	-0.000624	-1.631743e-04	
LotShape_Reg	0.002302	0.000000e+00	
LandContour_HLS	0.009553	8.059010e-03	
LandContour_Low	0.000639	-0.000000e+00	
LandContour_Lvl	0.012923	1.129460e-02	
Utilities_NoSeWa	-0.008475	-4.876407e-03	
LotConfig_CulDSac	0.010431	9.964450e-03	
LotConfig_FR2	-0.006903	-3.813893e-03	

	ScreenPorch	0.011493	1.109327e-02
PoolArea	0.103497	1.435153e-01	
MiscVal	0.004650	-0.000000e+00	
MoSold	0.000685	0.000000e+00	
Age	-0.053685	-5.381423e-02	

```
[166]: lasso_cols_removed = list (betas[betas['Lasso']==0].index)
print(lasso_cols_removed)

['MSSubClass_45', 'MSSubClass_50', 'MSSubClass_75', 'MSSubClass_80', 'MSSubClass_85', 'MSSubClass_120', 'MSSubClass_190', 'MSZoning_RM', 'LotShape_Reg', 'LandContour_Low', 'LandSlope_Sev', 'Neighborhood_Blueste', 'Neighborhood_BrDale', 'Neighborhood_CollGr', 'Neighborhood_Gilbert', 'Neighborhood_Sawyer_W', 'Condition2_Norm', 'Condition2_PosA', 'Condition2_RRan', 'BldgType_2fmCon', 'HouseStyle_2.5Fin', 'HouseStyle_2Story', 'HouseStyle_S1', 'OverallQual_6', 'OverallCond_6', 'RoofStyle_Hip', 'RoofMatl_Membran', 'Exterior1st_CementBd', 'Exterior1st_Plywood', 'Exterior1st_VinylSd', 'Exterior2nd_Brk_Cmn', 'Exterior2nd_CBlock', 'Exterior2nd_HdBoard', 'Exterior2nd_MetalSd', 'Exterior2nd_Stucco', 'Exterior2nd_VinylSd', 'Exterior2nd_Wd_Shng', 'MasVnrType_BrkFace', 'ExterQual_Gd', 'ExterCond_Gd', 'Foundation_CBlock', 'BsmtQual_Fa', 'BsmtQual_No', 'BsmtCond_Non', 'BsmtExposure_Non', 'BsmtFinType2_GLO', 'BsmtFinType2_Non', 'BsmtFinType2_Rec', 'BsmtFinType2_Umf', 'Heating_GasA', 'Electrical_FuseF', 'Electrical_Mix', 'KitchenQual_Fa', 'Functional_Mini', 'FireplaceQu_Gd', 'GarageType_Basment', 'GarageType_Builtin', 'GarageType_CarPort', 'GarageType_Detchd', 'GarageQual_Po', 'GarageQual_TA', 'GarageCond_Gd', 'GarageCond_TA', 'Fence_MnPrv', 'MiscFeature_Shed', 'SaleType_ConL1', 'SaleType_ConW', 'SaleType_WD', 'SaleCondition_Family', 'LotFrontage', 'BsmtUnfSF', '2ndFlrSF', 'LowQualFinSF', 'BsmtHalfBath', 'GarageYrBlt', 'MiscVal', 'MoSold']

[170]: lasso_cols_selected = list (betas[betas['Lasso']!=0].index)
print(lasso_cols_selected)

['MSSubClass_45', 'MSSubClass_50', 'MSSubClass_75', 'MSSubClass_80', 'MSSubClass_85', 'MSSubClass_120', 'MSSubClass_190', 'MSZoning_RM', 'LotShape_Reg', 'LandContour_Low', 'LandSlope_Sev', 'Neighborhood_Blueste', 'Neighborhood_BrDale', 'Neighborhood_CollGr', 'Neighborhood_Gilbert', 'Neighborhood_Sawyer_W', 'Condition2_Norm', 'Condition2_PosA', 'Condition2_RRan', 'BldgType_2fmCon', 'HouseStyle_2.5Fin', 'HouseStyle_2Story', 'HouseStyle_S1', 'OverallQual_6', 'OverallCond_6', 'RoofStyle_Hip', 'RoofMatl_Membran', 'Exterior1st_CementBd', 'Exterior1st_Plywood', 'Exterior1st_VinylSd', 'Exterior2nd_AsphShn', 'Exterior2nd_Brk_Cmn', 'Exterior2nd_CBlock', 'Exterior2nd_HdBoard', 'Exterior2nd_MetalSd', 'Exterior2nd_Stucco', 'Exterior2nd_VinylSd', 'Exterior2nd_Wd_Shng', 'MasVnrType_BrkFace', 'ExterQual_Gd', 'ExterCond_Gd', 'Foundation_CBlock', 'BsmtQual_Fa', 'BsmtQual_No']
```

Surprise Housing Case

```
[171]: print(len(lasso_cols_removed))
print(len(lasso_cols_selected))

80
80

[172]: betas['Ridge'].sort_values(ascending=False)[:10]

RoofMatl_CompShg    0.118084
PoolArea            0.103497
RoofMatl_Tar&Grv   0.087644
PoolQC_None         0.083895
RoofMatl_WdShngl   0.059597
MSZoning_RL         0.059178
GrLivArea           0.057481
MSZoning_RM         0.046035
RoofMatl_WdShake   0.045562
2ndFlrSF            0.039642
Name: Ridge, dtype: float64

[173]: ridge_coeffs = np.exp(betas['Ridge'])
ridge_coeffs.sort_values(ascending=False)[:10]

RoofMatl_Compshg    1.125339
PoolArea            1.109042
RoofMatl_Tar&Grv   1.091599
PoolQC_None         1.087515
```

Surprise Housing Case

```
[174]: betas['Lasso'].sort_values(ascending=False)[:10]

RoofMatl_WdShngl   1.061408
MSZoning_RL         1.060965
GrLivArea           1.059165
MSZoning_RM         1.047111
RoofMatl_WdShake   1.046616
2ndFlrSF            1.040438
Name: Lasso, dtype: float64

[175]: betas['Lasso'].sort_values(ascending=False)[:10]

PoolArea            0.143515
PoolQC_None         0.125196
GrLivArea           0.107108
RoofMatl_CompShg   0.038784
RoofMatl_Tar&Grv   0.036527
OverallQual_9        0.032135
OverallQual_8        0.031684
TotalBsmtSF          0.029316
RoofMatl_WdShngl   0.023278
Neighborhood_Crawfor 0.022705
Name: Lasso, dtype: float64

[176]: lasso_coeffs = np.exp(betas['Lasso'])
lasso_coeffs.sort_values(ascending=False)[:10]

PoolArea            1.154325
PoolQC_None         1.133371
GrLivArea           1.113054
RoofMatl_CompShg   1.039546
RoofMatl_Tar&Grv   1.037202
OverallQual_9        1.032657
OverallQual_8        1.032192
TotalBsmtSF          1.029750
```

THANK YOU

